




Damien Aicheh

Cloud Solution Architect

damienaicheh@microsoft.com



Format of this training


- Training for Terraform **beginners**
- Alternate between **theory and practice**
- **6 blocs of 4 hours** to focus on Terraform
- For the time of the training, you're blocked from customer and other internal calls
- There are **no stupid questions** so don't hesitate to ask
- At the end of this training, you will have a good understanding of the way to start a Terraform project

Prerequisites for each participant

- An Azure account with **contributor** rights
- Visual Studio Code with Terraform plugin:
 - <https://marketplace.visualstudio.com/items?itemName=hashicorp.terraform>
- In your local machine:
 - Terraform
 - Azure CLI
 - Docker (optional)
- Access to GitHub Actions / Ability to create repositories



GitHub Actions



HashiCorp Terraform

HashiCorp

- HashiCorp is based in San Francisco and was founded in **2012**
- HashiCorp provides a suite of **open-source tools** intended to support development and deployment of large-scale service-oriented software installations.
- Each tool is aimed at **specific stages in the life cycle of a software** application, with a focus on automation

HashiCorp current products

INFRASTRUCTURE



Terraform

Infrastructure as code



Packer

Machine images

SECURITY



Vault

Identity-based security



Boundary

Secure remote access

NETWORKING



Consul

Multi-cloud service
networking

APPLICATIONS



Nomad

Workload orchestration



Waypoint

App deployment
workflows



Vagrant

Environment workflows

Terraform

- HashiCorp Terraform enables you to safely and predictably **create, change, and improve infrastructure**.
- **Codifies APIs** into **declarative configuration** files that can be shared among team members, treated as code, edited, reviewed, and versioned.
- **Abstraction layer** over common public/private cloud provisioning/automation languages
- Declarative Config Files are written in **HCL** (HashiCorp Configuration Language)
- **Single binary** without complex installation
- **Cross platform**, can be run on Windows, Linux, or MacOS
- Written in Go with **plugin-oriented architecture**

Terraform

- What Terraform is **NOT** write once, deploy anywhere
- **Providers** are the libraries for the different vendors
- Providers have **different settings**
- Terraform enables deployment in Azure, AWS and Google using **different** providers, configured with **different settings**
- Terraform use the **same language** but **not** the **same resource**:
 - Azure: Storage Account
 - AWS: S3
 - GCP: Cloud Storage



Terraform

- Terraform enables **management of multi cloud** environments using a single toolset and syntax
- Broad **community**, very **responsive**
- **Frequent releases:** approximately every two weeks
- **Extensible**
- Open source / **no vendor** lock-in
- Microsoft is heavily investing in making **most Azure services supported**
- Microsoft **add Terraform support** on there products, like Azure Dev Cli (azd)

Terraform: Caveats

- Very **new, or advanced Azure functionality** may not be natively available in Terraform Provider
- Workarounds do exist: the **azurerm_template_deployment** resource can be used to **wrap a native ARM Template**, but this has its own limitations (it doesn't support resource deletions)
- **Unencrypted secrets** in state file
- You have to manage the **state** of your platform

Licencing

- On August 10, 2023, HashiCorp announced a change of license for its products
- Terraform being open source under the **MPL v2 license**, it was to move under a non-open source **BSL v1.1 license**
- Starting from the version **1.6**
- At this point it impacts only the company **building products based on** the Terraform technology
- Companies that deploy resources using Terraform today can **continue** to do it **for free**
- Initiative with Opentofu: <https://github.com/opentofu/manifesto>

Hands On Lab!

2 options:

- Install everything on your machine
- Use DevContainers inside Visual Studio Code and Docker:

<https://marketplace.visualstudio.com/items?itemName=ms-vscode-remote.remote-containers>

https://github.com/damienaicheh/terraform_tooling



Check your configuration

Just run in your favorite terminal:

For Terraform:



```
terraform --version
```

For Azure CLI:



```
az --version
```

Set up your Azure account



```
# Login to Azure  
az login
```

```
# Show your account  
az account show
```

```
# Select your subscription  
az account set --subscription <your-subscription-id>
```

HCL Language

HCL (HashiCorp Configuration Language)

One language and Cross Platform

Human and machine **readable**

*.TF files, UTF-8 encoded

Blocks containing key/value pairs

HCL **is not JSON**

- Less nesting
- Less brackets

Supports inline comments

#, //, /* */

Case sensitive

Naming convention: **snack case**

```
resource "azurerm_resource_group" "example" {
  name      = "example-resources"
  location  = "West Europe"
}

// Comments here...
resource "azurerm_storage_account" "example" {
  name                        = "storageaccountname"
  resource_group_name        = azurerm_resource_group.example.name
  location                   = azurerm_resource_group.example.location
  account_tier               = "Standard"
  account_replication_type   = "GRS"

  tags = {
    environment = "staging"
  }
}
```

ARM Template vs Terraform HCL

Terraform HCL is about 50-70 % more compact than ARM.

ARM JSON



```
"name": "[concat(parameters('PilotServerName'), '-vm')]",
```

Terraform HCL



```
name = "${var.PilotServerName}-vm"
```

ARM Template vs Terraform HCL

ARM Template	Terraform
JSON	HCL
Parameters	Variables
Variables	Local variables
Resources	Resources
Functions	Functions
Nested templates	Modules
Explicit dependency	Automatic dependency
Refer by reference or resourceid	Refer by resource or data source

Terraform variables, locals and outputs

	Variables	Locals	Outputs
What	Parameters	Local variables	Results
Example	Environment name	Resource group name based on combination of variables, operations, prefix and suffix	Resource Identifier, Resource Url
How to reference it?	var.<name>	local.<name>	module.<module name>.<output name>

How to declare variables, locals and outputs



```
variable "resource_group_name_suffix" {  
  type      = string  
  description = "The resource group name suffix"  
}
```



```
locals {  
  resource_name_suffix = format("%s-%s", var.environment, var.owner)  
}
```



```
output "app_service_url" {  
  value = azurerm_app_service.this.default_site_hostname  
  description = "The URL of the App Service"  
}
```

Optional arguments: Description

Always a string



```
variable "resource_group_name_suffix" {  
  type      = string  
  description = "The resource group name suffix"  
}
```



```
output "app_service_url" {  
  value      = azurerm_app_service.this.azurerm_app_service  
  description = "The URL of the App Service"  
}
```

Optional arguments: sensitive

Boolean (true/false)

Suppressing Values in CLI Output

```
variable "my_service_password" {  
  type      = string  
  description = "My service password"  
  sensitive = true  
}
```



Secrets must be stored in Key Vault directly

Secrets are stored unencrypted in the state

```
output "db_password" {  
  value      = aws_db_instance.db.password  
  description = "The password for logging in to the database."  
  sensitive  = true  
}
```

Optional arguments: depends_on

Explicit dependency


Does not exist for variables

```
output "default_site_hostname" {  
    value      = azurerm_app_service.this.default_site_hostname  
    description = "The default site hostname."  
  
    depends_on = [  
        azurerm_app_service.this  
    ]  
}
```


Optional arguments: default

Default value means that the variable become optional

Does not exist for outputs



```
variable "resource_group_name_suffix" {  
    type      = string  
    default   = "01"  
    description = "The resource group name suffix"  
}
```

Optional arguments: type

Types is optional but recommended to restrict the values like:

- string
- number
- bool
- list<Type>
- set(<Type>)
- map(<Type>)
- object({<ATTR NAME> = <TYPE>, ... })
- tuple([<TYPE>, ...])

Does not exist for outputs

Optional arguments: type

Category	Type	Description	Example: Declare	Example: Access Value
Any	Any	Any type / no constraint		var.myVar
Primitive Type	string	Double quoted, with common escape sequences: \n, \t, \\, \" Multi-line strings „heredoc“ format.	\"\"multi\"\"\\nline\" <<EOT multiline heredoc string EOF	var.myVar
	number	whole number, as well as fractions	12 3.3475	var.myVar
	bool	true / false		var.myVar
Complex Types	list(<TYPE>)	sequence of values, 0-based index	[\"us-west-1a\", \"us-west-1c\"]	var.myVar[0]
	set(<TYPE>)	Collection of unique values		var.myVar[0]
	tuple([<TYPE>, ...])	Similar to list, but each element may have different types	[\"us-west-1a\", 12]	var.myVar[0]
	map(<TYPE>)	group of values identified by named labels	{name = \"Mabel\", city = \"Munich\"}	var.myVar.name var.myVar[\"name\"]
	object({<ATTR NAME> = <TYPE>, ... })	Similar to map, but attributes may have different types	{name = \"Mabel\", age = 52}	var.myVar.name var.myVar[\"name\"]

Optional arguments: validation

Advanced validation rules like regex or conditions


Does not exist for outputs

```
variable "resource_group_name_suffix" {  
  type      = string  
  default   = "01"  
  description = "The resource group name suffix"  
  validation {  
    condition     = can(regex("[0-9]{2}", var.resource_group_name_suffix))  
    error_message = "The resource group name suffix value must be two digits."  
  }  
}
```

Optional arguments: nullable

Boolean (true/false), whether null is allowed or not for the variable, default is true

Does not exist for outputs



```
variable "environment" {  
  type      = string  
  nullable  = false  
}
```

Terraform Expressions

Category	Syntax	Description
Operators	<ol style="list-style-type: none">1. <code>!, -</code> (multiplication by -1)2. <code>*, /, %</code>3. <code>+, -</code> (subtraction)4. <code>>, >=, <, <=</code>5. <code>==, !=</code>6. <code>&&</code>7. <code> </code>	<p>Ordered operators == requires same type / no implicit conversion</p> <p>In general: Terraform is case sensitive => <code>"a" != "A"</code> => Also case for function names matters!</p>
Conditional expression	<p><code>condition ? true_val : false_val</code> e.g. <code>var.a != "" ? var.a : "default-a"</code></p>	<p>Implicit conversion, if results don't have same type => <code>string + number</code> results in string</p>
For expression	<p><code>[for k, v in var.map : length(k) + length(v)]</code></p>	<p>Variable for key k is optional. In case of map and object, it contains the key. For list, set and tuple, it contains the index number.</p>

Terraform Functions

Category	Functions - see official docs for full list
Numeric	<ul style="list-style-type: none">• Functions with 1 or 2 param: abs, ceil, floor, log, parseint, pow, signum• max and min: numeric functions expecting 1..n arguments<ul style="list-style-type: none">• Syntax: max(1, 3, 4)• Or expand list using "...": max([1, 3, 4]...)• sum: collection function, only expecting list: sum([1, 3, 4])
String	chomp, format, formatlist, join, lower, regex, regexall, replace, split, strrev, substr, title, trim, upper
String / Collection	length
Collections	<ul style="list-style-type: none">• Checks: alltrue, anytrue, contains,• Returning collections: chunklist, coalescelist, compact, concat, distinct, flatten, keys, range, reverse, slice, values• Returning primitive types: coalesce, index, lookup,• Returning maps or objects: merge, transpose, zipmap• Returning sets: setintersection, setproduct, setsubtract, setunion
Encoding	base64gzip, csvdecode, jsondecode, jsonencode, textdecodebase64, textencodebase64, urlencode, yamldecode, yamlencode
Date and time	formatdate, timeadd, timestamp
Hash and Crypto	md5, sha256, sha512, rsadecrypt, uuid
Type Conversion	tobool, tolist, tomap, tonumber, toset, toString, try

Terraform Workflow and commands

Terraform Workflow: init



init


- **Initializing** working directory
- **Download** provider version
- **Analyzing** modules
- **Backend** initialization

Terraform Workflow: init

- Before running *terraform init*, **connect to Azure** using az CLI to obtain an **authorization token**, and select the **subscription** you wish to deploy to.
- Run the **init** command in the directory that contains the ***.tf*** files to be deployed.
- It determines which **providers** to download based on the .tf files in the current directory.
- Provider files are stored within the current directory in a hidden directory ***".terraform"***

```
PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL

PS C:\Repos\Terraform_Demo\BigDayIn\resourcegroup> terraform init

Initializing provider plugins...
- Checking for available provider plugins on https://releases.hashicorp.com...
- Downloading plugin for provider "azurerm" (1.27.1)... 

The following providers do not have any version constraints in configuration,
so the latest version was installed.

To prevent automatic upgrades to new major versions that may contain breaking
changes, it is recommended to add version = "..." constraints to the
corresponding provider blocks in configuration, with the constraint strings
suggested below.


* provider.azurerm: version = "~> 1.27"

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
```

Terraform Workflow: plan



- **Initializing** working directory
 - **Download** provider version
 - **Analyzing** modules
 - **Backend** initialization
- **Compares** code changes to current state
 - Creates a list of instruction to **migrate the state to the desired changes**
 - Creates a **plan file** with the list of instruction (no execution)

Terraform Workflow: plan

- **Plan** command provide a **preview of the actions**
- It's an opportunity to **sanity check** the deployment.
- Resource **additions** will be indicated by a **+** symbol, and **highlighted in green**
- Resource **modifications** will be indicated by a **~** symbol and **highlighted in orange**
- Resource **deletions** will be indicated by a **-** symbol and **highlighted in red**
- **Colours** only shown **if supported by terminal**
- Where required, specify the **-var** and **-var-file** parameters to pass in variable values to the **plan** command.

```
PS C:\Repos\Terraform_Demo\BigDayIn\resourcegroup> terraform plan
Refreshing Terraform state in-memory prior to plan...
The refreshed state will be used to calculate this plan, but will not be
persisted to local or remote state storage.

-----

An execution plan has been generated and is shown below.
Resource actions are indicated with the following symbols:
  + create

Terraform will perform the following actions:


+ azurerm_resource_group.rg
  id:          <computed>
  location:    "australiaeast"
  name:        "DEVOPStestresourcegroup"
  tags.%:      <computed>

Plan: 1 to add, 0 to change, 0 to destroy.

-----

Note: You didn't specify an "-out" parameter to save this plan, so Terraform
can't guarantee that exactly these actions will be performed if
"terraform apply" is subsequently run.
```

Terraform Workflow: apply



- **Initializing** working directory
- **Download** provider version
- **Analyzing** modules
- **Backend** initialization

- **Compares** code changes to current state
- Creates a list of instruction to **migrate the state to the desired changes**
- Creates a **plan file** with the list of instruction (no execution)

- Actual infrastructure: **Applies** the set of instruction created in the plan file

Terraform Workflow: apply

- The **apply** command is used to execute the terraform deployment and add, change and delete resources.
- For automated deployments the **-auto-approve** parameter is used to skip the confirmation step.
- Where required, specify the **-var** and **-var-file** parameters to pass in variable values to the **apply** command.
- This command **can take some time** to run for large deployments. Continuous output will be provided to assist in monitoring progress.

```
PS C:\Repos\Terraform_Demo\BigDayIn\resourcegroup> terraform apply

An execution plan has been generated and is shown below.
Resource actions are indicated with the following symbols:
  + create

Terraform will perform the following actions:

+ azurerm_resource_group.rg
  id:          <computed>
  location:    "australiaeast"
  name:        "DEVOPStestresourcegroup"
  tags.%:      <computed>

Plan: 1 to add, 0 to change, 0 to destroy.


Do you want to perform these actions?
  Terraform will perform the actions described above.
  Only 'yes' will be accepted to approve.

Enter a value: yes

azurerm_resource_group.rg: Creating...
  location: "" => "australiaeast"
  name:      "" => "DEVOPStestresourcegroup"
  tags.%:    "" => "<computed>"
azurerm_resource_group.rg: Creation complete after 3s (ID: /subscriptions/85b

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.
```

Terraform Workflow: destroy



- **Initializing** working directory
- **Download** provider version
- **Analyzing** modules
- **Backend** initialization

- **Compares** code changes to current state
- Creates a list of instruction to **migrate the state to the desired changes**
- Creates a **plan file** with the list of instruction (no execution)

- Actual infrastructure: **Applies** the set of instruction created in the plan file

- Actual infrastructure: **Remove** all configured resources

Terraform Workflow: destroy

- The **destroy** command will remove all resources specified in the deployment .tf files.
- When executed, a **plan is printed** to show the resources that will be destroyed, and you are prompted for confirmation.
- For automated deployments the **-auto-approve** parameter is used to skip the confirmation step.
- Where required, specify the **-var** and **-var-file** parameters to pass in variable values to the **destroy** command.
- This command **can take some time** to run for large deployments. Continuous output will be provided to assist in monitoring progress.

```
PS C:\Repos\Terraform_Demo\BigDayIn\resourcegroup> terraform destroy
azurerm_resource_group.rg: Refreshing state... (ID: /subscriptions/85bff22c-722c-487d-b4

An execution plan has been generated and is shown below.
Resource actions are indicated with the following symbols:
  - destroy

Terraform will perform the following actions:

  - azurerm_resource_group.rg

Plan: 0 to add, 0 to change, 1 to destroy.

Do you really want to destroy all resources?
  Terraform will destroy all your managed infrastructure, as shown above.
  There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

azurerm_resource_group.rg: Destroying... (ID: /subscriptions/85bff22c-722c-487d-b41d-...
azurerm_resource_group.rg: Still destroying... (ID: /subscriptions/85bff22c-722c-487d-b4
azurerm_resource_group.rg: Still destroying... (ID: /subscriptions/85bff22c-722c-487d-b4
azurerm_resource_group.rg: Still destroying... (ID: /subscriptions/85bff22c-722c-487d-b4
azurerm_resource_group.rg: Still destroying... (ID: /subscriptions/85bff22c-722c-487d-b4
azurerm_resource_group.rg: Still destroying... (ID: /subscriptions/85bff22c-722c-487d-b4
azurerm_resource_group.rg: Destruction complete after 50s

Destroy complete! Resources: 1 destroyed.
```


Terraform commands

Command	Description
terraform init	Initialize the context of the project (providers, modules, backends...)
terraform fmt -recursive	Format the tf files in a standard manner
terraform validate	Validate the syntax of the tf files BUT there is no guarantee that the deployment will succeed
terraform plan	Define the potential creation, updates, deletion based on the current tfstate of the architecture
terraform apply	Apply the changes based on the plan.out if provided
terraform destroy	Destroy all the resources defined in the tfstate

Hands On Lab!

HCL: HashiCorp Configuration Language

- Play with variable file provided (main.tf)
- Use **terraform init** to initialize the project
- To test the code just run:

terraform plan

terraform apply

- The instructions are on the next slide


HCL: HashiCorp Configuration Language

- Output the **application_name** with only the 3 first characters
- Output the **letters** array in **one single string** with **dashes** to separate each letter
- Create a variable **string** called **domain** and when you output it make sure it's **lowercase**
- Return the **node count** in a dedicated output
- Return the **node size** in a dedicated output
- Output a resource group name which will be a **concatenation** of the official abbreviation "rg" and the **application name** and **domain** formatted previously. Dashes must separate each part.
E.g. rg-<application name>-<domain>

Resources & Data Sources

Resources

- Most **important element** in the Terraform language.
- Each resource block **describes** one or more **infrastructure** objects
- Syntax:



```
resource "[Resource-Type]" "Resource-Name" {  
    [Optional parameters]  
}
```

- <https://www.terraform.io/language/resources>

Resources

- Syntax example for an Azure resource group:



```
resource "azurerm_resource_group" "this" {  
  name      = var.resource_group_name  
  location = var.location  
}
```


- Syntax to access to values:



```
azurerm_resource_group.this.location
```

Data Sources

- Use information defined **outside of Terraform**, for instance Azure
- Use information defined by **another** separate Terraform **configuration** (state)
- Syntax:



```
data "[Data-Source-Type]" "Data-Name" {  
  [Optional parameters]  
}
```

- <https://www.terraform.io/language/data-sources>

Data Sources

- Syntax example for data source to declare the access the configuration of the AzureRM provider:



```
data "azurerm_client_config" "current" {}
```

- Syntax to access the values, for instance the subscription id:



```
data.azurerm_client_config.current.subscription_id
```

Terraform State

Terraform State: Overview


- Terraform **stores the state** about your managed infrastructure and configuration **in a file**
- The state is used to **map real world resources to the configuration**
- Terraform **uses this state to create plans** and make changes to your infrastructure
- **Sensitive information** is **stored unencrypted** within state
- **terraform show** can be used to display managed resources in the state

Terraform State: Local vs. Remote

- The state is stored by default in a **local file** named **terraform.tfstate**
- It's configured using backend block in the provider resource
- Local State – for **non-prod tests only**
- Remote State – Should be **stored remotely** (e.g. blob) to enable **security and collaboration**
- Before each deployment, Terraform must be able to **obtain a lock on the state file**.
- Once the **deployment is complete**, the **lock is released**.
- **Azure Storage Accounts** support **file locking** which prevents multiple people from simultaneously executing a Terraform deployment, **security**, **version control** and **disaster recovery** capabilities for your state files.

Terraform State: Remote

- Terraform has his own cloud: **Terraform Cloud Remote State Management**
- **Azure Storage Accounts** support **file locking** which prevents multiple people from simultaneously executing a Terraform deployment, **security**, **version control** and **disaster recovery** capabilities for your state files.

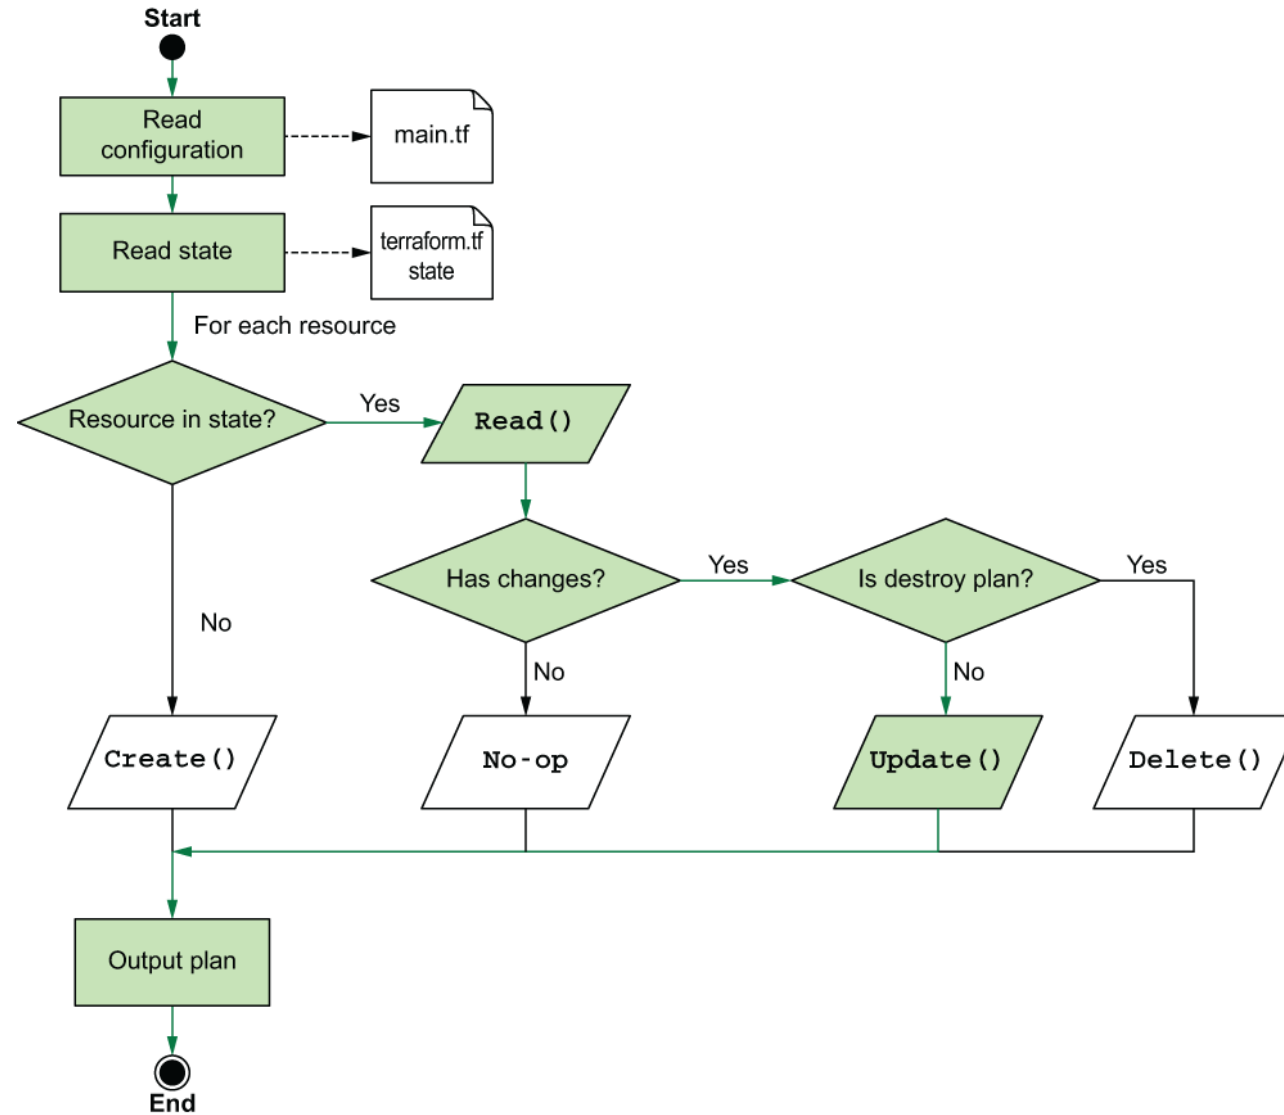


- If a state is **lock**, you will have to connect to Azure Storage Account and click “Break lease”

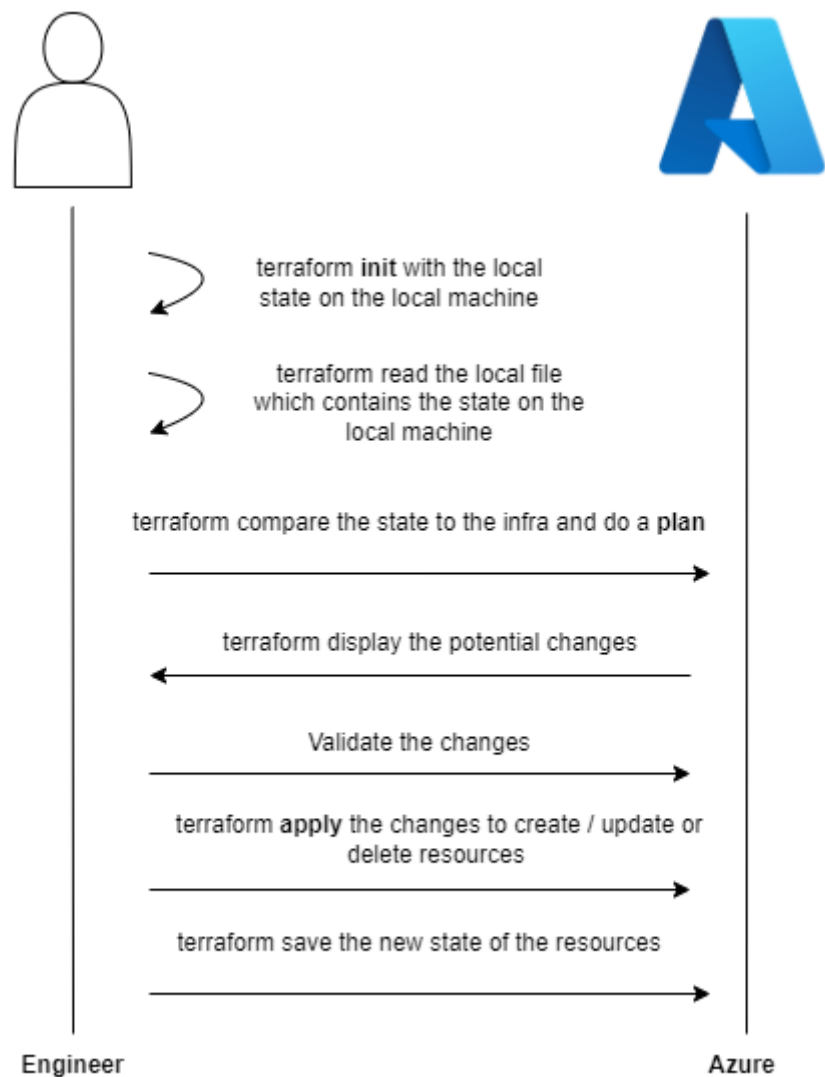
terraform.tfstate ...

Discard Download Refresh Delete Change tier Acquire lease **Break lease** Give feedback

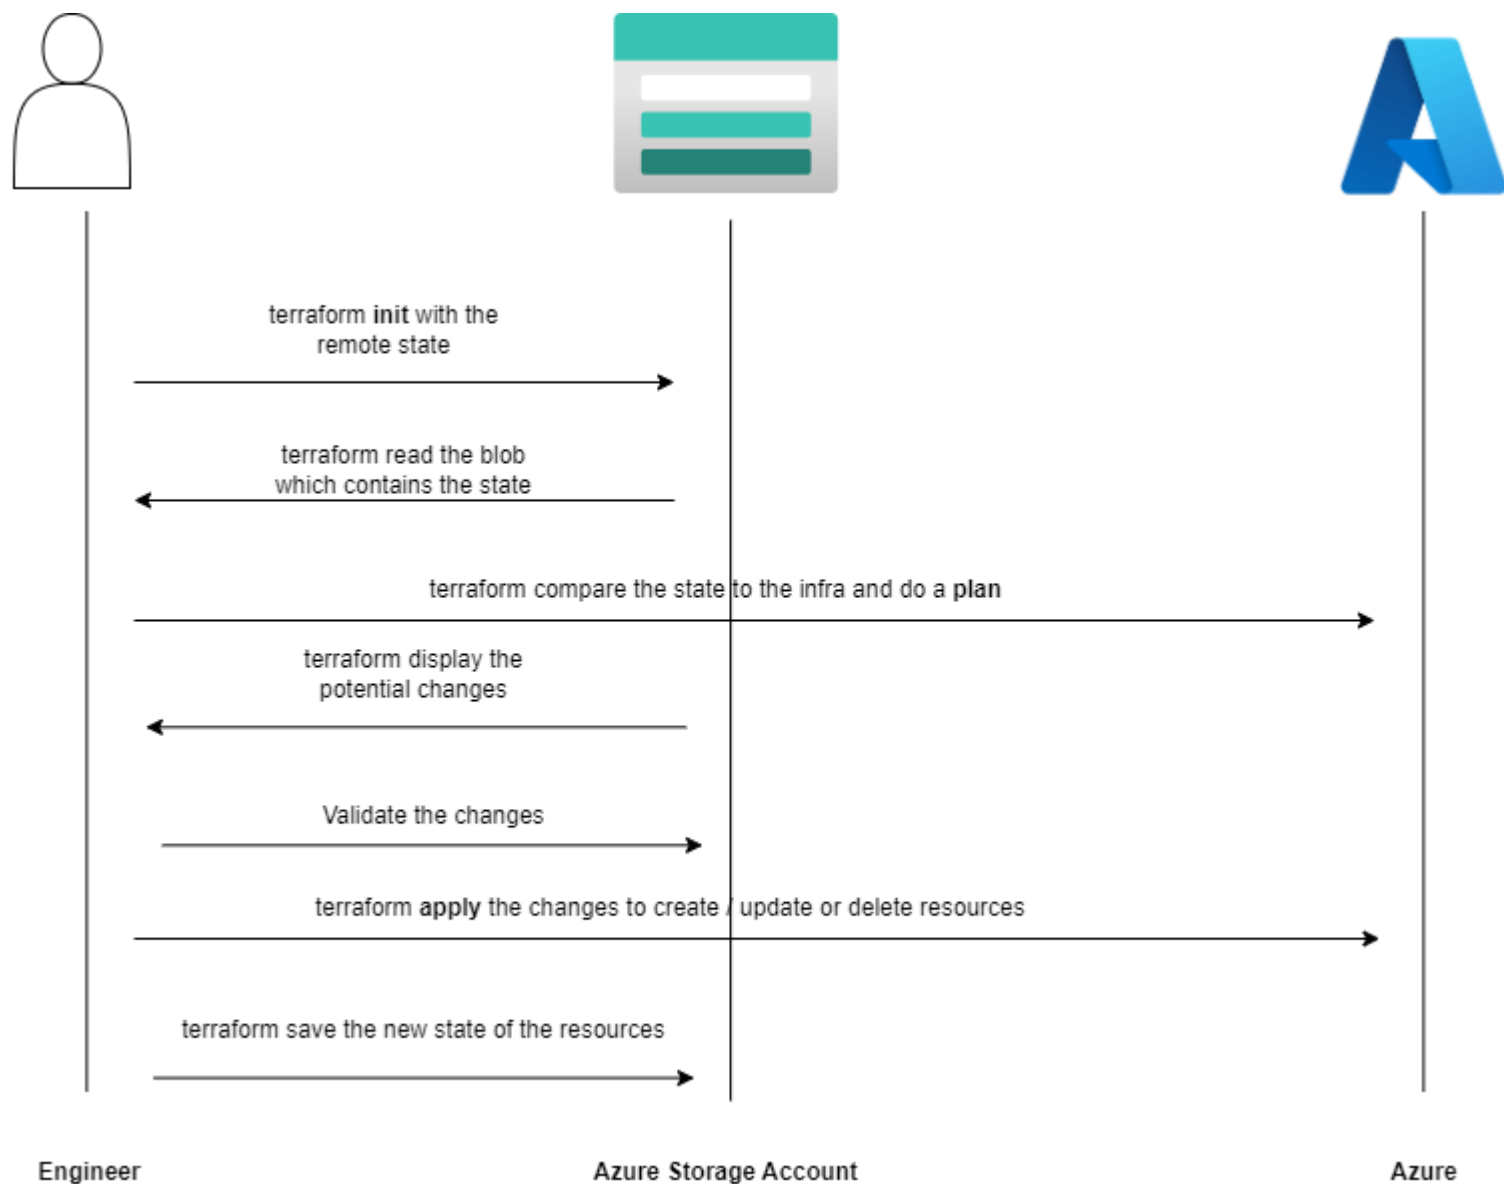
How the plan is generated



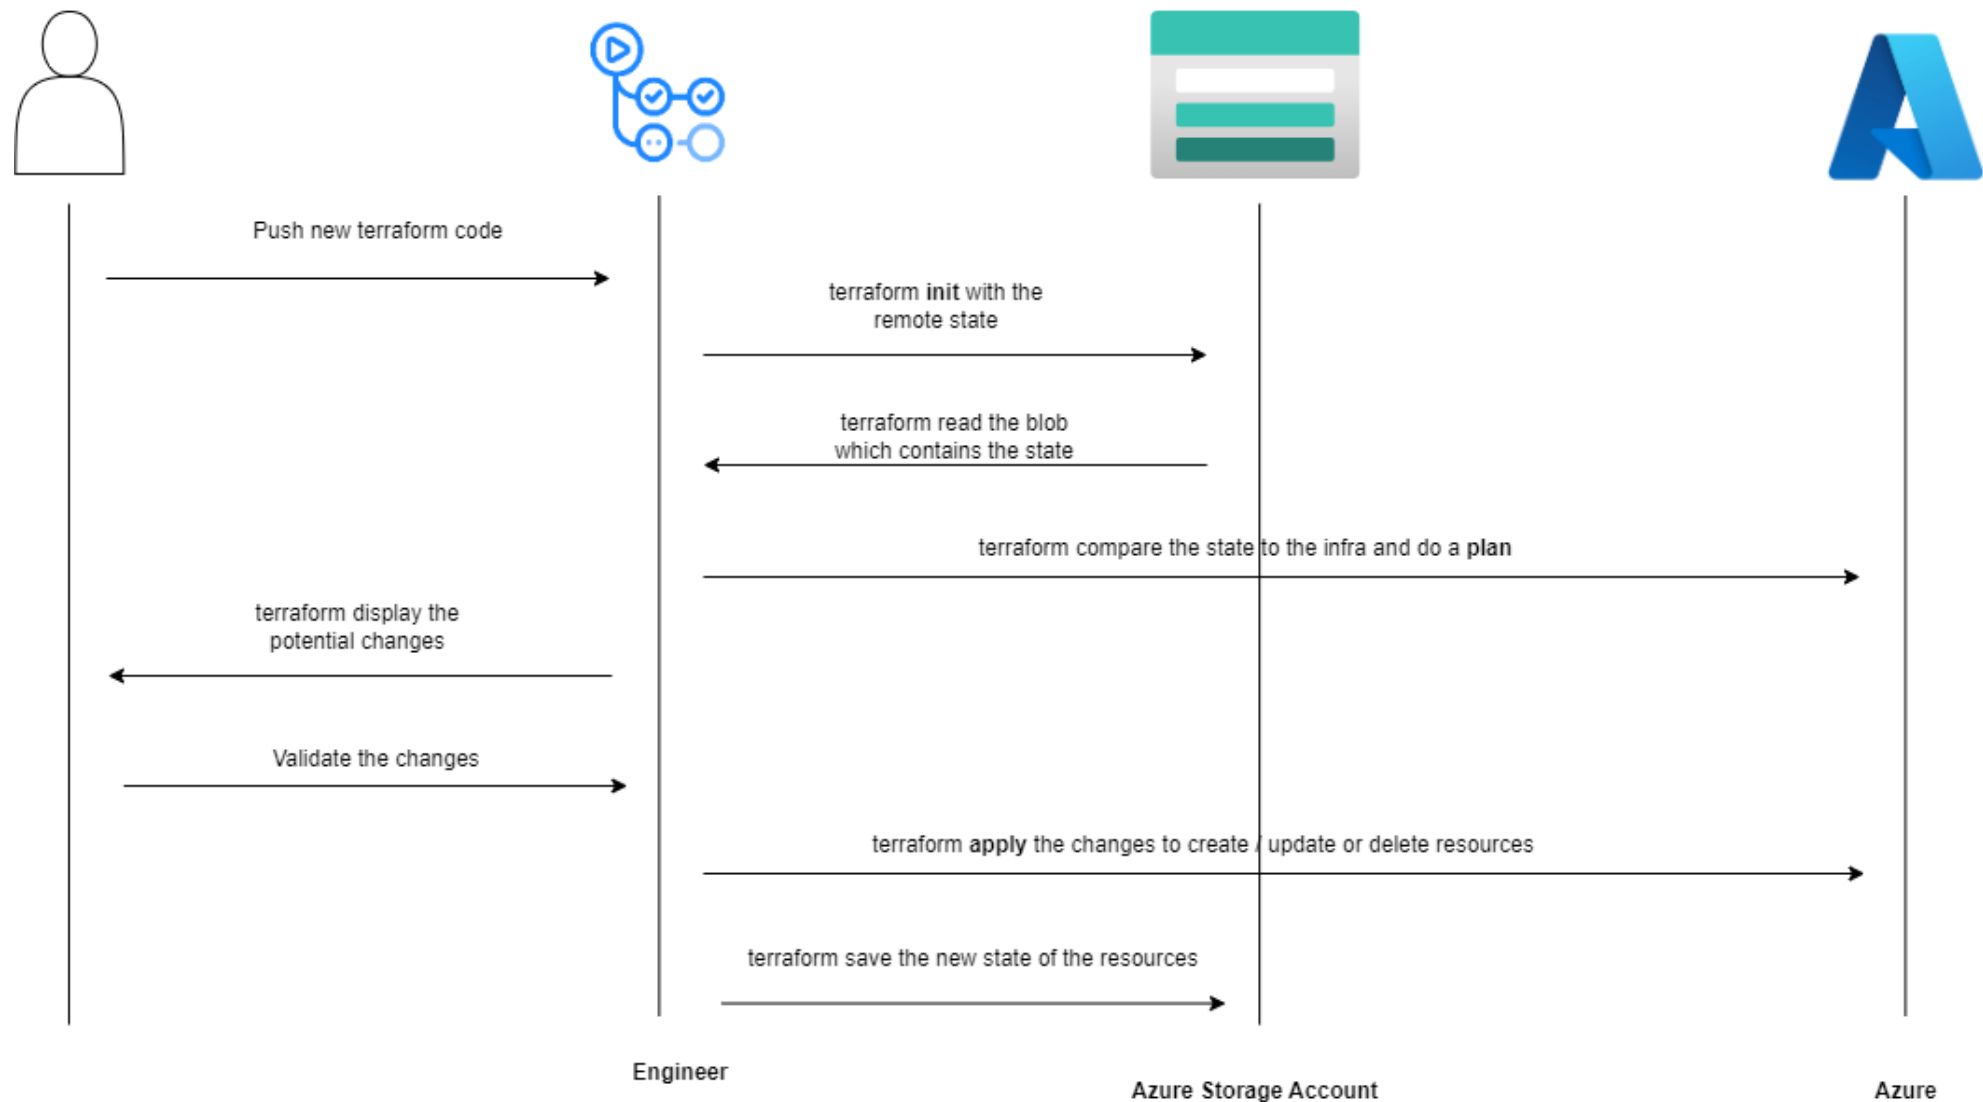
Terraform State: Locally



Terraform State: Remotely



Terraform State: DevOps



Don't touch the User Interface



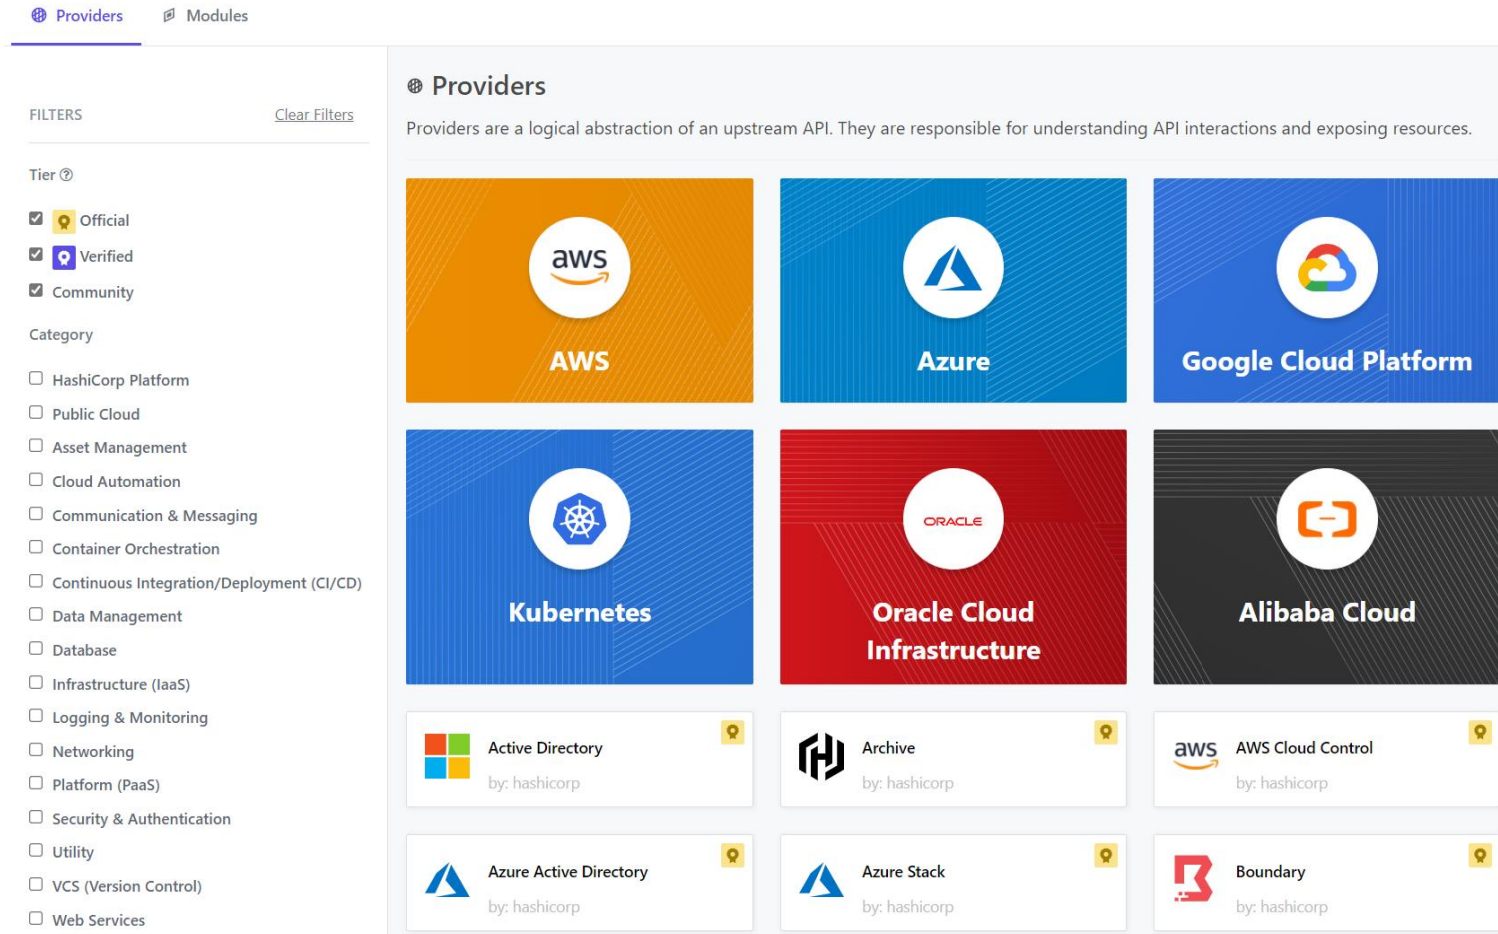
- Keep the User Interface in **ReadOnly** mode
- If you modify a resource managed by Terraform you can end up with a **broken state**
- You will have to import each resource **manually** to the state, which can be really long and complex
- **terraform import** can be used to explicitly map actual resources to configurations.



```
terraform import -var-file ./env.dev.tfvars azurerm_resource_group.this  
/subscriptions/733df85b-e409-4b3d-8ca2-abb90bb4a2c2/resourceGroups/my-  
resource-group
```

Providers

Terraform Provider: Available ones



The screenshot shows the Terraform Registry's Providers page. On the left, there's a sidebar with filters. The 'Providers' tab is selected. The filters include 'Tier' (Official, Verified, Community) and 'Category' (HashiCorp Platform, Public Cloud, Asset Management, Cloud Automation, Communication & Messaging, Container Orchestration, Continuous Integration/Deployment (CI/CD), Data Management, Database, Infrastructure (IaaS), Logging & Monitoring, Networking, Platform (PaaS), Security & Authentication, Utility, VCS (Version Control), Web Services). The main content area is titled 'Providers' and includes a description: 'Providers are a logical abstraction of an upstream API. They are responsible for understanding API interactions and exposing resources.' Below this, there's a grid of provider cards. The first row features AWS, Azure, and Google Cloud Platform. The second row features Kubernetes, Oracle Cloud Infrastructure, and Alibaba Cloud. The third row features Active Directory, Archive, and AWS Cloud Control. The fourth row features Azure Active Directory, Azure Stack, and Boundary. Each card displays the provider's logo and name, and many have a yellow 'Official' badge.

Providers

Providers are a logical abstraction of an upstream API. They are responsible for understanding API interactions and exposing resources.

Filters:

Tier

- ☒ Official
- ☒ Verified
- ☒ Community

Category

- ☐ HashiCorp Platform
- ☐ Public Cloud
- ☐ Asset Management
- ☐ Cloud Automation
- ☐ Communication & Messaging
- ☐ Container Orchestration
- ☐ Continuous Integration/Deployment (CI/CD)
- ☐ Data Management
- ☐ Database
- ☐ Infrastructure (IaaS)
- ☐ Logging & Monitoring
- ☐ Networking
- ☐ Platform (PaaS)
- ☐ Security & Authentication
- ☐ Utility
- ☐ VCS (Version Control)
- ☐ Web Services


Providers Grid:

- AWS**
- Azure**
- Google Cloud Platform**
- Kubernetes**
- Oracle Cloud Infrastructure**
- Alibaba Cloud**
- Active Directory** (by: hashicorp)
- Archive** (by: hashicorp)
- AWS Cloud Control** (by: hashicorp)
- Azure Active Directory** (by: hashicorp)
- Azure Stack** (by: hashicorp)
- Boundary** (by: hashicorp)

- <https://registry.terraform.io/browse/providers>


Terraform Providers

- The terraform core program **requires at least one** provider to build anything.
- You can **manually configure** which **version(s)** of a provider you would like to use.
- If you leave this option out, Terraform will **default to the latest** available version of the provider.



```
provider "azurerm" {  
  version = "=3.58.0"  
}
```

Configure azurearm provider




```
terraform {  
  required_providers {  
    azurearm = {  
      source  = "hashicorp/azurearm"  
      version = "3.64.0"  
    }  
  }  
}  
  
backend "local" { } # backend "azurearm" {  
}  
  
provider "azurearm" {  
  # Configuration options  
}
```

Authentication with AzureRm

With Service Principal Name (SPN)

- Use a user account for **development purpose**
- Use a Service Principal Name (SPN) or Az CLI for **deployment purpose**
- The SPN should have **Contributor** role



```
az ad sp create-for-rbac --name <service_principal_name> --role Contributor --scopes /subscriptions/<subscription_id>
```

<https://learn.microsoft.com/en-us/azure/developer/terraform/authenticate-to-azure?tabs=bash>

Service Principal Name directly in the code

- You can connect Terraform to Azure directly in the **code using the provider**
- SPN credentials in your code is a **bad practice**
- You will probably **forget to remove** them before pushing the code
- This is a **security** issue

```
provider "azurerm" {  
  features {}  
  
  subscription_id = "<azure_subscription_id>"  
  tenant_id       = "<azure_subscription_tenant_id>"  
  client_id       = "<service_principal_appid>"  
  client_secret    = "<service_principal_password>"  
}
```

NOT RECOMMENDED

Service Principal Name with environment variable


- You can connect Terraform to Azure using **environment variables**
- SPN credentials will be **only on your machine**
- Your code doesn't have **any credentials**
- Credential lifecycle is a problem with SPN, **you need to manage it**

```
export ARM_SUBSCRIPTION_ID="<azure_subscription_id>"  
export ARM_TENANT_ID="<azure_subscription_tenant_id>"  
export ARM_CLIENT_ID="<service_principal_appid>"  
export ARM_CLIENT_SECRET="<service_principal_password>"
```

<https://learn.microsoft.com/en-us/azure/developer/terraform/authenticate-to-azure?tabs=bash>

With Az CLI

- You can connect Terraform to Azure using **Az CLI directly**:



```
# Login to Azure
az login

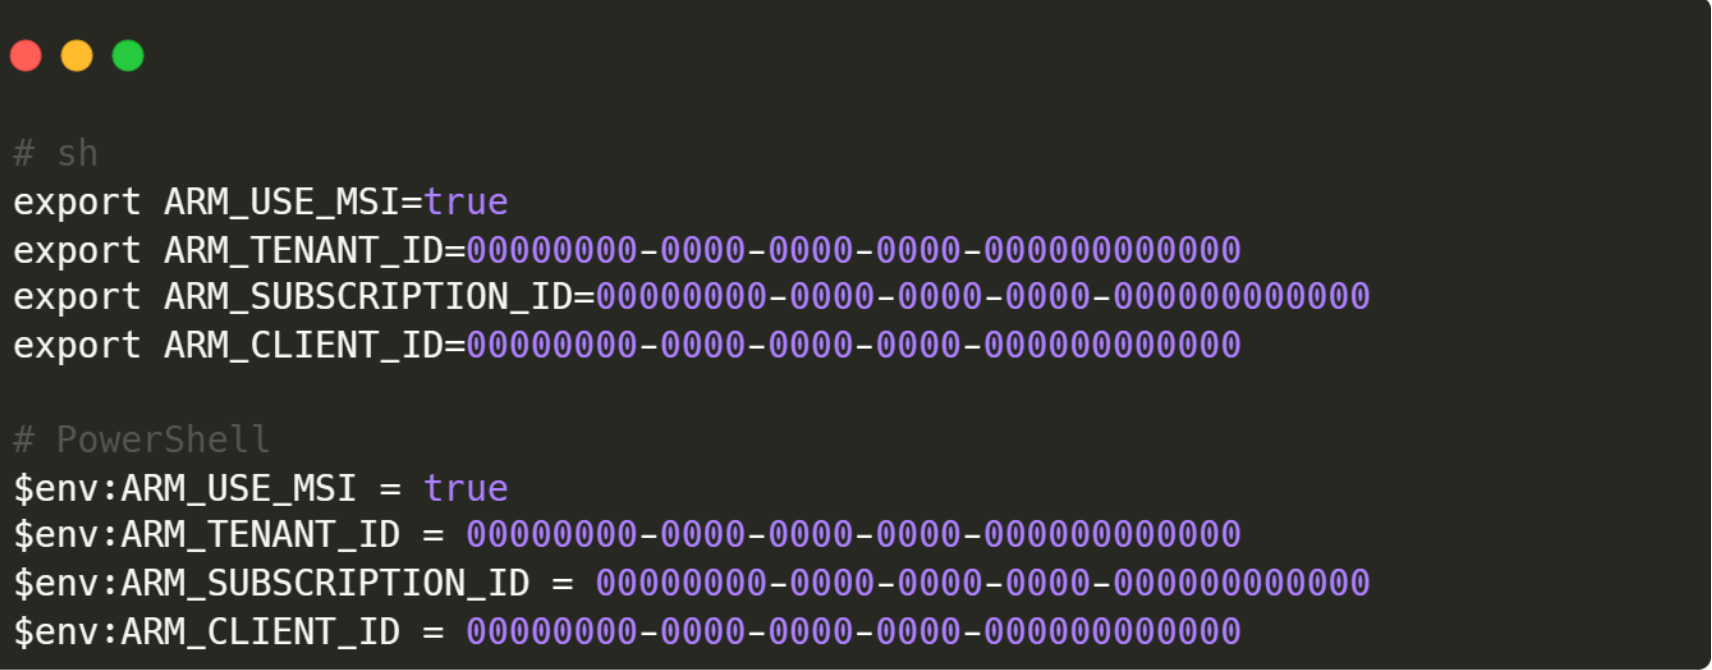
# Show your account
az account show

# Select your subscription
az account set --subscription <your-subscription-id>
```

- Use it for **development purpose only**

With Managed Identity

- Use managed identity declared in **Azure Active Directory**
- The token to authenticate to Azure is managed automatically by Azure Active Directory. You don't need **to monitor the end of life** of your token



```
# sh
export ARM_USE_MSI=true
export ARM_TENANT_ID=00000000-0000-0000-0000-000000000000
export ARM_SUBSCRIPTION_ID=00000000-0000-0000-0000-000000000000
export ARM_CLIENT_ID=00000000-0000-0000-0000-000000000000

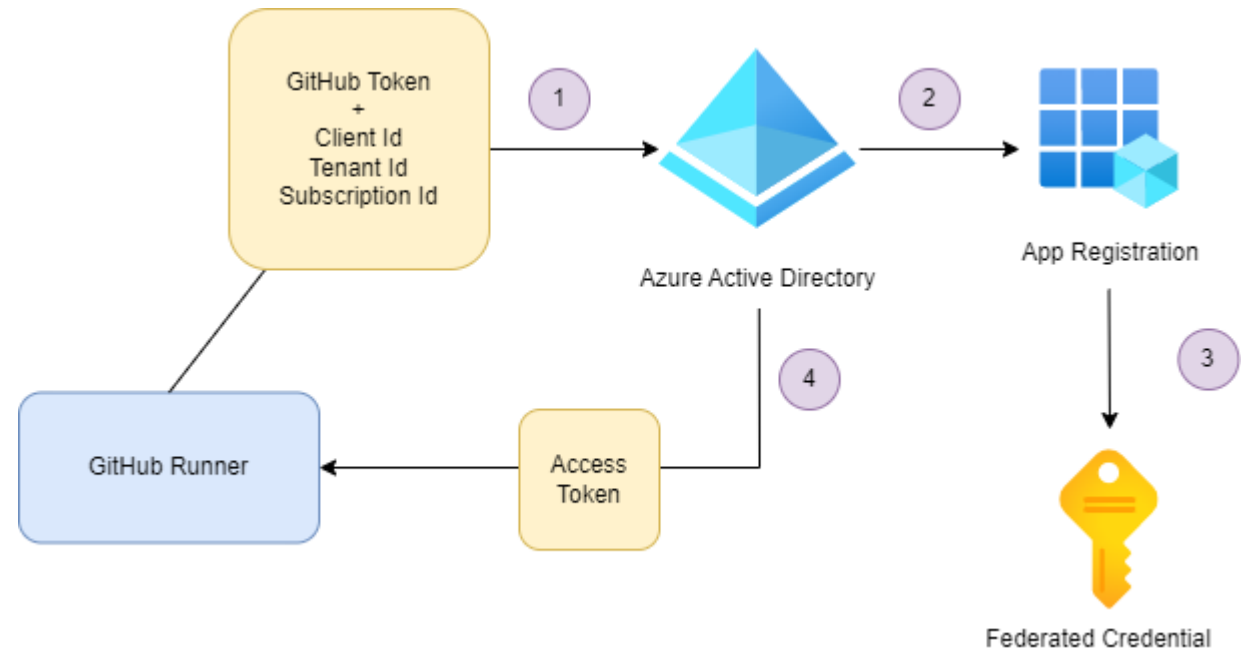
# PowerShell
$env:ARM_USE_MSI = true
$env:ARM_TENANT_ID = 00000000-0000-0000-0000-000000000000
$env:ARM_SUBSCRIPTION_ID = 00000000-0000-0000-0000-000000000000
$env:ARM_CLIENT_ID = 00000000-0000-0000-0000-000000000000
```

With Open ID Connect (OIDC)

- Azure **manage** the credentials for you
- Generate credentials with **short lifetime**
- Avoid managing manually a **Service Principal Name** (SPN)
- Just need to:
 - Create an **App Registration**
 - Create **federated credentials** to give access to the repository

OIDC: How it works?

- 1) GitHub Runner asks with a GitHub **token** **and an Azure AD identity** for an AAD token. Then token contains issuer, audience and subject.
- 2) Based on the token AAD search the associated identity in the **App Registration**
- 3) The App Registration has the **federated credential** for this repository
- 4) AAD return an access token, so the GitHub Runner can **modify Azure Resources**



Hands On Lab!

Connect to Azure Rm

- Create a new folder called **digital_infra_state**
- Create a provider.tf file and use use **Azure Rm**
<https://registry.terraform.io/providers/hashicorp/azurerm/latest/docs>
- Use **the latest version** of **Azure Rm**
- Use **local state**
- Run **terraform init** to validate that everything is **correctly initialized**

Connect to Azure Rm

Initializing the backend...

Initializing provider plugins...

- Reusing previous version of hashicorp/azurerm from the dependency lock file
- Using previously-installed hashicorp/azurerm v3.59.0

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see any changes that are required for your infrastructure. All Terraform commands should now work.

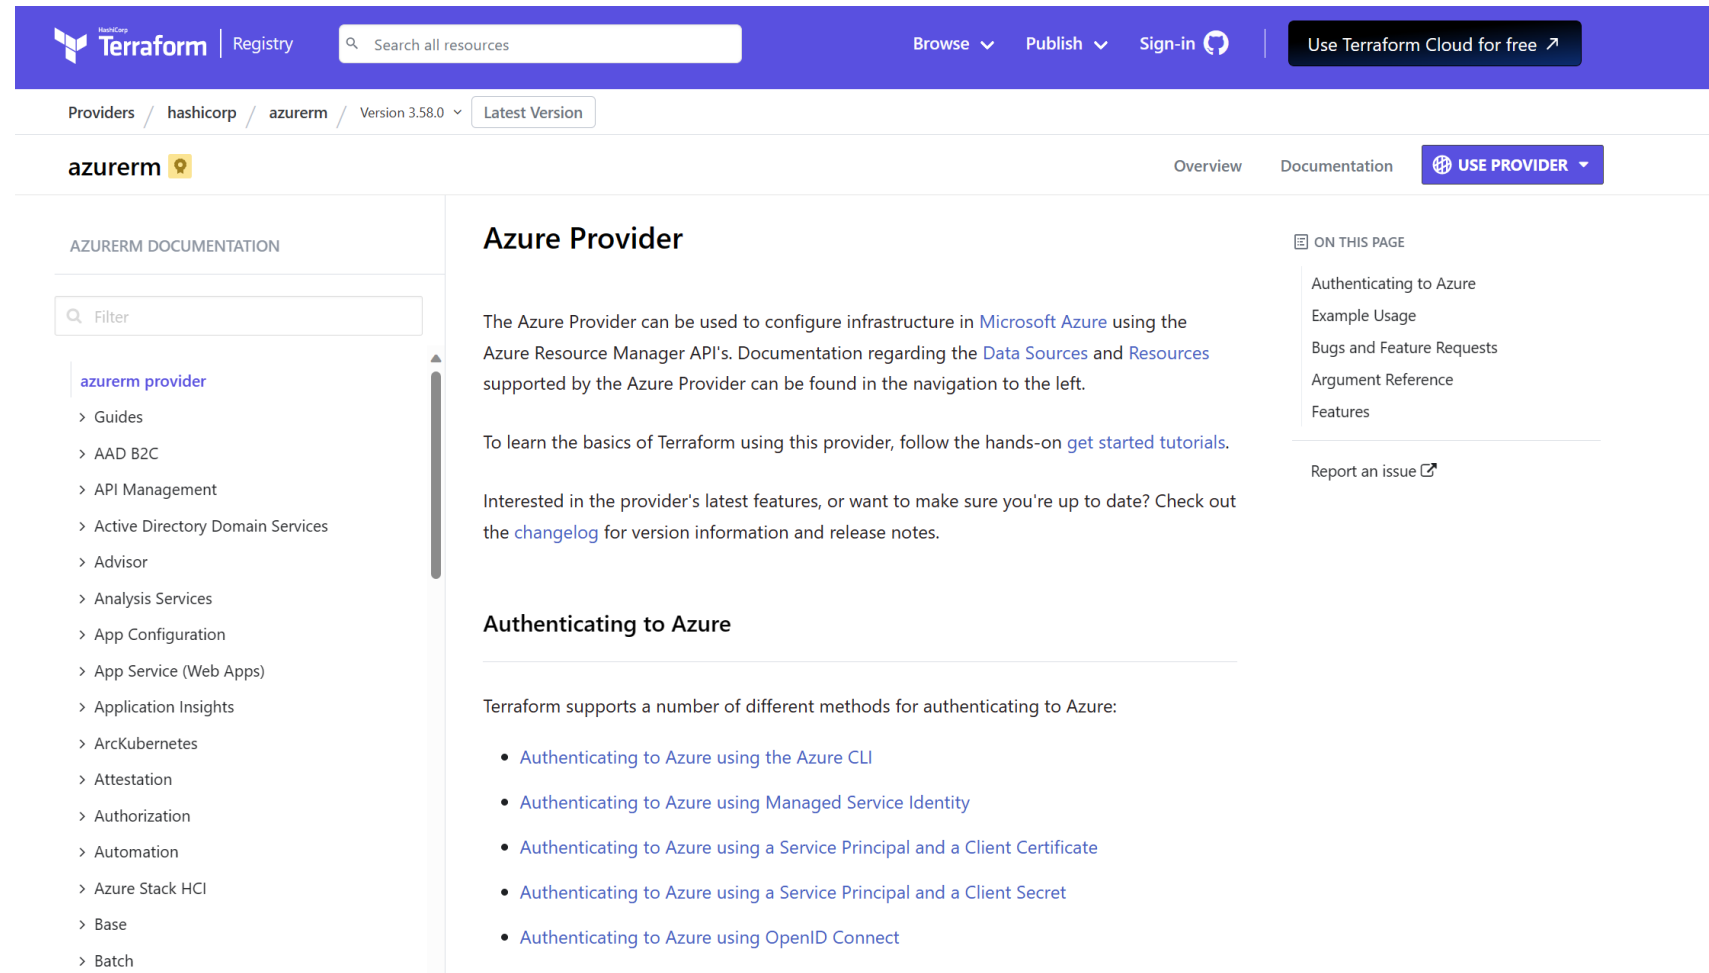
If you ever set or change modules or backend configuration for Terraform, rerun this command to reinitialize your working directory. If you forget, other commands will detect it and remind you to do so if necessary.

Documentation usage: Overview

Documentation


- Really good documentation
- The search bar can be a bit capricious
- Use your favorite search engine to find the resource. For instance, type: Resource group Terraform

<https://registry.terraform.io/providers/hashicorp/azurerm/latest/docs>



The screenshot shows the Terraform Registry page for the Azure Provider. The header is purple with the Terraform logo, a search bar, and navigation links like 'Browse', 'Publish', and 'Sign-in'. Below the header, the breadcrumb trail shows 'Providers / hashicorp / azurerm / Version 3.58.0 / Latest Version'. The main content area is titled 'azurerm' and has tabs for 'Overview' and 'Documentation'. The 'Documentation' tab is active, showing the 'AZURERM DOCUMENTATION' section with a search bar and a list of links. The 'Azure Provider' section explains that the provider is used to configure infrastructure in Microsoft Azure using the Azure Resource Manager API's. It also mentions that documentation regarding the Data Sources and Resources supported by the Azure Provider can be found in the navigation to the left. A 'USE PROVIDER' button is visible in the top right. On the right side, there is a 'ON THIS PAGE' section with links to 'Authenticating to Azure', 'Example Usage', 'Bugs and Feature Requests', 'Argument Reference', and 'Features'. At the bottom right, there is a 'Report an issue' link.

Providers / **hashicorp** / **azurerm** / Version 3.58.0 / Latest Version

azurerm 

Overview Documentation **USE PROVIDER**

AZURERM DOCUMENTATION

Filter

- azurerm provider
- > Guides
- > AAD B2C
- > API Management
- > Active Directory Domain Services
- > Advisor
- > Analysis Services
- > App Configuration
- > App Service (Web Apps)
- > Application Insights
- > ArcKubernetes
- > Attestation
- > Authorization
- > Automation
- > Azure Stack HCI
- > Base
- > Batch

Azure Provider

The Azure Provider can be used to configure infrastructure in [Microsoft Azure](#) using the Azure Resource Manager API's. Documentation regarding the [Data Sources](#) and [Resources](#) supported by the Azure Provider can be found in the navigation to the left.

To learn the basics of Terraform using this provider, follow the hands-on [get started tutorials](#).

Interested in the provider's latest features, or want to make sure you're up to date? Check out the [changelog](#) for version information and release notes.

Authenticating to Azure

Terraform supports a number of different methods for authenticating to Azure:

- [Authenticating to Azure using the Azure CLI](#)
- [Authenticating to Azure using Managed Service Identity](#)
- [Authenticating to Azure using a Service Principal and a Client Certificate](#)
- [Authenticating to Azure using a Service Principal and a Client Secret](#)
- [Authenticating to Azure using OpenID Connect](#)

ON THIS PAGE

- [Authenticating to Azure](#)
- [Example Usage](#)
- [Bugs and Feature Requests](#)
- [Argument Reference](#)
- [Features](#)

[Report an issue](#)

Resource sections

- Always cut into several parts:
 - **Basic examples** to copy-paste
 - **Argument** References which are the available **properties** of the resource
 - **Attributes** References which are the available **outputs** of the resource
 - **Timeouts** properties that can be **optionally** specified
 - **Import** a resource to the **current state**

Example

HashiCorp

Terraform

Registry

Search all resources

Browse

Publish

Sign-in

Use Terraform Cloud for free

Providers

hashicorp

azurerm

Version 3.58.0

Latest Version

azurerm

Overview

Documentation

USE PROVIDER

AZURERM DOCUMENTATION

azurerm storage account

54 matching results

App Service (Web Apps)

Resources

azurerm_source_control_token

Data Sources

azurerm_source_control_token

Automation

Resources

azurerm_automation_account

Data Sources

azurerm_automation_account

Batch

Resources

azurerm batch account

azurerm_storage_account

Manages an Azure Storage Account.

Example Usage

```
resource "azurerm_resource_group" "example" {
  name      = "example-resources"
  location  = "West Europe"
}

resource "azurerm_storage_account" "example" {
  name                        = "storageaccountname"
  resource_group_name        = azurerm_resource_group.example.name
  location                   = azurerm_resource_group.example.location
  account_tier                = "Standard"
  account_replication_type    = "GRS"

  tags = {
    environment = "staging"
  }
}
```

ON THIS PAGE

Example Usage

Example Usage with Network Rules

Argument Reference

Attributes Reference

Timeouts

Import

Report an issue

©Microsoft Corporation
Azure

Arguments

Argument Reference

The following arguments are supported:

- `name` - (Required) Specifies the name of the storage account. Only lowercase Alphanumeric characters allowed. Changing this forces a new resource to be created. This must be unique across the entire Azure service, not just within the resource group.
- `resource_group_name` - (Required) The name of the resource group in which to create the storage account. Changing this forces a new resource to be created.
- `location` - (Required) Specifies the supported Azure location where the resource exists. Changing this forces a new resource to be created.
- `account_kind` - (Optional) Defines the Kind of account. Valid options are `BlobStorage`, `BlockBlobStorage`, `FileStorage`, `Storage` and `StorageV2`. Defaults to `StorageV2`.


NOTE:

Changing the `account_kind` value from `Storage` to `StorageV2` will not trigger a force new on the storage account, it will only upgrade the existing storage account from `Storage` to `StorageV2` keeping the existing storage account in place.

- `account_tier` - (Required) Defines the Tier to use for this storage account. Valid options are `Standard` and `Premium`. For `BlockBlobStorage` and `FileStorage` accounts only `Premium` is valid. Changing this forces a new resource to be created.

ON THIS PAGE

- Example Usage
- Example Usage with Network Rules
- [Argument Reference](#)
- Attributes Reference
- Timeouts
- Import

[Report an issue](#) 

Attributes


Attributes Reference

In addition to the Arguments listed above - the following Attributes are exported:

- `id` - The ID of the Storage Account.
- `primary_location` - The primary location of the storage account.
- `secondary_location` - The secondary location of the storage account.
- `primary_blob_endpoint` - The endpoint URL for blob storage in the primary location.
- `primary_blob_host` - The hostname with port if applicable for blob storage in the primary location.
- `secondary_blob_endpoint` - The endpoint URL for blob storage in the secondary location.
- `secondary_blob_host` - The hostname with port if applicable for blob storage in the secondary location.
- `primary_queue_endpoint` - The endpoint URL for queue storage in the primary location.

ON THIS PAGE

- Example Usage
- Example Usage with Network Rules
- Argument Reference
- [Attributes Reference](#)
- Timeouts
- Import

[Report an issue](#) 

Hands On Lab!

Create a resource group with terraform local state

- Back to your **digital_infra_state** folder
- Create a variable.tf file to define :
 - An environment which accept only "dev", "stag", "prod" with "dev" as default value
 - A domain name with default value to "state"
 - A location with a default value set to "westeurope"
 - A resource suffix with a default value
- Create a locals.tf file to define two **resource suffix** which are a concatenation of:
 - The environment
 - The first 2 letters of the location in lowercase
 - The first 3 letters of the domain in lowercase
 - The resource suffix


The resource suffixes should be in this format:

With dashes (kebabcase): <env>-<domain>-<location>-<resource-suffix>

Without dashes (lowercase): <env><domain><location><resource-suffix>

Create a resource group with terraform local state

- Create a rg.tf file to define your resource group. Naming convention is: **rg-<your-resource-suffix-kebabcase>**
- Create an sto.tf file to define a basic **storage account**:
The naming convention is **st<your-resource-suffix-lowercase>**
Account tier: **Standard**
Account replication type: **LRS**
- Inside the Storage Account create **one private container** called "**states**"
- Add tags "Domain" and "Environment" to all your resources.
- Use **local state**
- **Deploy** to Azure



Backend definition

Backend configuration

- The backend configuration for Azure Rm is a **Storage Account**
- The Storage Account contains a **container** with a **blob** which is the **tfstate**
- Which came first, the chicken or the egg?

Connect to the Azure Rm backend in code

- Useful for testing purpose **BUT**:
- Avoid using this approach because your code must be agnostic to the targeted environment

NOT RECOMMENDED

```
terraform {  
  backend "azurerm" {  
    resource_group_name = "digital_infra_state"  
    storage_account_name = "stostate01"  
    container_name      = "tfstates"  
    key                  = "prod.terraform.tfstate"  
  }  
}
```

Connect to the Azure Rm backend with backend config

- Use the **-backend-config** option to pass the Azure Rm settings
- This can be done for CI / CD
- **Use environment variables**

BEST PRACTICE

```
terraform init \  
  -backend-config="resource_group_name=digital_infra_state" \  
  -backend-config="storage_account_name=stostate01" \  
  -backend-config="container_name=tfstates" \  
  -backend-config="key=prod.terraform.tfstate"
```

Several ways to use environment variables

Pass variables to terraform commands


- Pass **-var** options on the command line



```
terraform apply -var="image_id=abc123"
```



Pass variables to terraform commands

- Pass a file as options with **-var-file**



```
terraform apply -var-file="testing.tfvars"
```

- **.tfvars** list all variables to set for the deployment

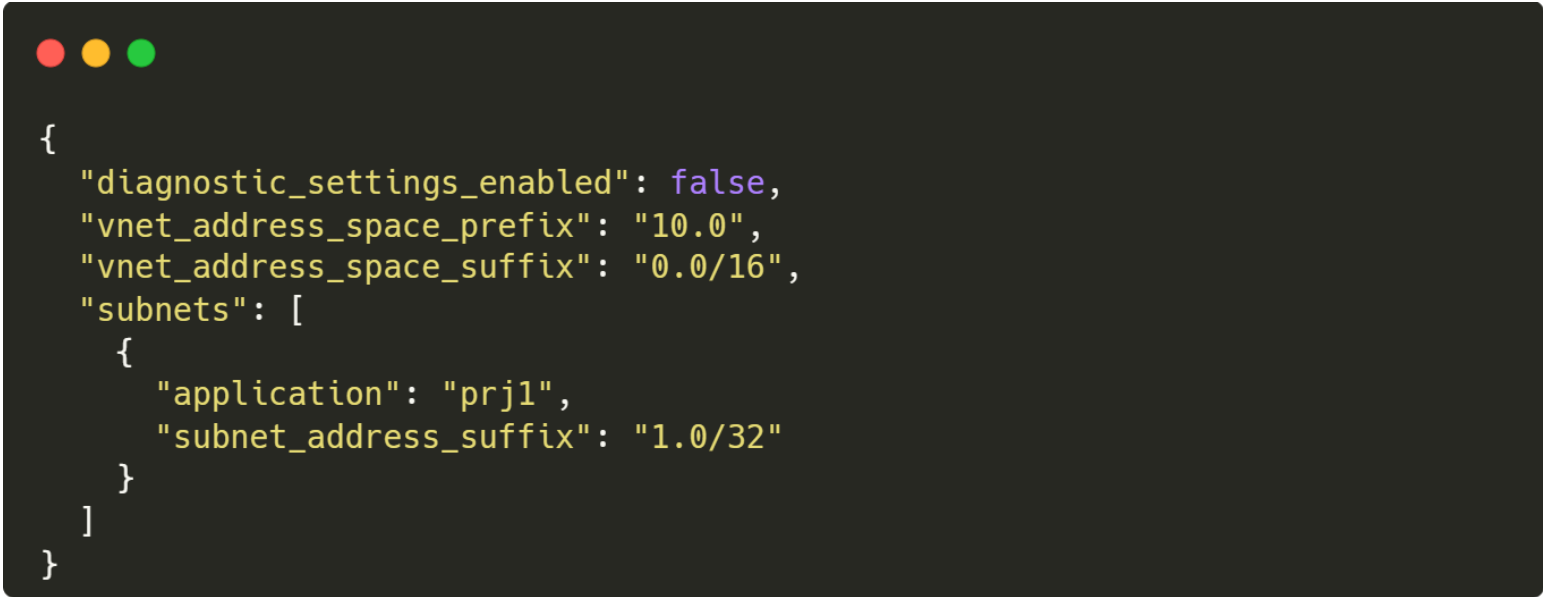


```
diagnostic_settings_enabled = false
vnet_address_space_prefix = "10.0"
vnet_address_space_suffix = "0.0/16"
subnets = [
  {
    "application" = "prj1"
    "subnet_address_suffix" = "1.0/32"
  }
]
```

Automatic load of variable files if present

Terraform also automatically loads a number of variable definitions files if they are present:

- Files named exactly **terraform.tfvars** or **terraform.tfvars.json**.
- Any files with names ending in **.auto.tfvars** or **.auto.tfvars.json**.



```
{
  "diagnostic_settings_enabled": false,
  "vnet_address_space_prefix": "10.0",
  "vnet_address_space_suffix": "0.0/16",
  "subnets": [
    {
      "application": "prj1",
      "subnet_address_suffix": "1.0/32"
    }
  ]
}
```

Other ways to pass variables

- Environment variables ("TF_VAR_[variable]", e.g. "TF_VAR_rg_id" for "rg_id")
- Default Config - default value in **variables.tf**
- User manual entry - if not specified, prompt the user for entry



```
export TF_VAR_image_id=ami-abc123  
terraform plan
```

Order of precedence of variables

- Once you have some variables defined, you can **set and override** them in **different ways**.
- Level of precedence for each method (1=highest, 7=lowest):
 1. **-var** and **-var-file** options on the command line
 2. Any ***.auto.tfvars** or ***.auto.tfvars.json** files, processed in lexical order of their filenames.
 3. The **terraform.tfvars.json** file, if present.
 4. The **terraform.tfvars** file, if present.
 5. Environment variables ("**TF_VAR_[variable]**", e.g. "**TF_VAR_rg_id**" for "**rg_id**")
 6. Default Config - default value in **variables.tf**
 7. User **manual entry** - if not specified, prompt the user for entry

Hands On Lab!

Create a Vnet and Subnet

- Create a new folder called **digital_infra_network**
- Based on the same model of the previous lab:
 - Create a **variable.tf**
 - Create a **locals.tf**
 - Create a **provider.tf**
 - Create an **rg.tf** for a new resource group
 - The domain will be the "**network**", so find a diminutive to it
- Use **remote state** and **store it** inside the Storage Account containers "states" created in the previous lab in this path: **network/terraform.state**

Create a Vnet and Subnet

- Create a **vnet.tf** file to define a VNet with an address range of: 10.0.0.0/16
- Create a **subnets.tf** file to define 2 subnets:
 - PaasSubnet: 10.0.0.0/24
 - AppServiceSubnet: 10.0.1.0/28 with a **delegation** to "Microsoft.Web/serverFarms" and **action** to "Microsoft.Network/virtualNetworks/subnets/action"
- Expose the VNet and Subnet Ids in an **outputs.tf** file
- Create a folder called **env-vars** for your environment variables and one folder for the **dev** environment. Create an **env.tfvars** file inside the **dev** folder to specify the values for the IP ranges for the VNet and the Subnets.
- Deploy only the **dev** environment

Meta-Arguments

Meta-Arguments: depends_on

- **Explicitly** specifying a dependency
- Necessary when a resource or module relies on some other resource's behavior but doesn't access to any of that resource's data in its arguments.

```
module "apim" {  
  source    = "./modules/api_management"  
  rg_name   = azurerm_resource_group.apifunc_rg.name  
  location  = azurerm_resource_group.apifunc_rg.location  
}
```

```
module "cdh_api_endpoints" {  
  source      = "./modules/cdh_api_endpoints"  
  rg_name     = azurerm_resource_group.apif_rg.name  
  apim_name   = "bmw-cdh-core-api"  
  api_management_name = var.api_apim_name  
  api_management_backend_name = "cdh-api"  
  depends_on = [  
    module.apim  
  ]  
}
```

Meta-Arguments: lifecycle

- Control lifecycle of a resource using arguments within lifecycle block:
create_before_destroy: bool, default false, if destroy is required, it is done after creating the new resource

prevent_destroy: bool, default false, plans resulting in a destroy will fail

ignore_changes: List of feature names (string), to be ignored, if different
Common scenarios to ignore later changes of initial values

default secrets in key vault

login of a sql server

app settings of a web app

- Only allows literal elements, no variables / expressions

```
resource "azurerm_function_app" "this" {  
  name = var.function_app_name  
  # ...  
  
  lifecycle {  
    ignore_changes = [app_settings]  
  }  
}
```

Meta-Arguments: provider

- Select non-default **provider** to be used
- Example: **multiple AzureRm provider** for different subscriptions

```
# Provider for current subscription
provider "azurerm" {
  features {}
}

# Provider for non-default Subscription
provider "azurerm" {
  alias                = "resource_subscription"
  subscription_id      = var.resource_subscription_id
  skip_provider_registration = true
  features {}
}
```

```
# RG in Default-Subscription
resource "azurerm_resource_group" "this" {
  name      = var.api_keyvault_resourcegroup
  location  = var.location
}


# RG in non-default subscription
resource "azurerm_resource_group" "this" {
  provider = azurerm.resource_subscription
  name      = var.api_keyvault_resourcegroup
  location  = var.location
}
```

Meta-Arguments: count

- Enables creation of **0..n similar objects**
- Count cannot be used simultaneously with meta-argument for_each.

Usage example:

- Zero-Based index number: count.index
- Individual instance: <TYPE>.<NAME>[<INDEX>] or module.<NAME>[<INDEX>]
- In the example below **Count = 0** means **skip creation**



```
resource "azurerm_monitor_metric_alert" "this" {  
  count          = var.add_default_alerts ? 1 : 0  
  name           = "Failed-Requests-Alert-${var.function_app_name}"  
  resource_group_name = var.rg_name  
  # ...  
}
```

Meta-Arguments: for_each

- Enables creation of **0..n similar objects**
- for_each cannot be used simultaneously with meta-argument count
- Use for_each instead of count for loop, this will avoid to have terraform state lost

Usage example:

- Current object: each.key (map key / set member), each.value (map value)
- Individual instance: <TYPE>.<NAME>[<KEY>] or module.<NAME>[<KEY>]

```
resource "azurerm_resource_group" "this" {  
  for_each = {  
    a_group      = "eastus"  
    another_group = "westus2"  
  }  
  name      = each.key  
  location = each.value  
}
```

Use for_each for the loops

- Use **for_each** instead of **count** for looping over a list of element.
- This will allow terraform to **keep tracking** existing resources and not recreating them



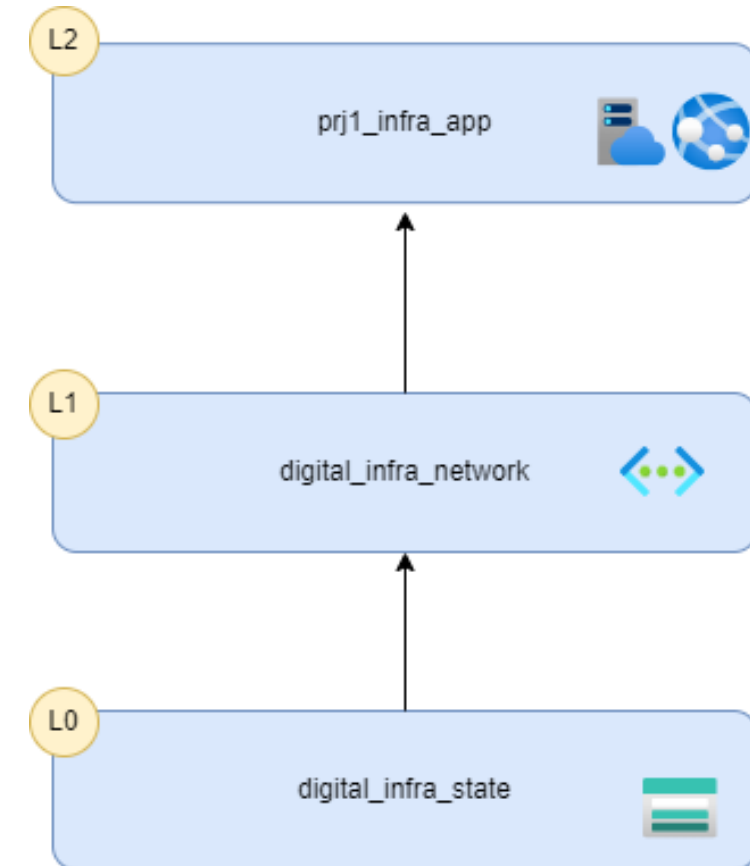

Hands On Lab!

Create an App Service

- Create a new folder called **prj1_infra_app**
- Based on the same model of the previous lab:
 - Create a **variable.tf**
 - Create a **locals.tf**
 - Create a **provider.tf**
 - Create an **rg.tf** for a new resource group
 - The domain will be the "**prj1**" and add a new variable for the application name and call it whatever you want but limit the characters to 3.
- Use **remote state** and **store it** inside the Storage Account containers "states" created in the previous lab in this path: **prj1/<application name>/terraform.state**

Create an App Service

- The goal of this lab is to consume the state network state and activate the VNet integration and private endpoint for the App Service.



Create an App Service

- Create an **app.tf** and define:
 - An App Service Plan with Basic SKU
 - An App Service for Linux
- Create a **data.tf** file and use the `terraform_remote_state` resource to connect to the network state with the correct credential methodology:
- <https://developer.hashicorp.com/terraform/language/settings/backends/azurerm>
- To be able to consume the state from another project you have to give the **Storage Blob Data Owner** role on the storage account to manage the states for Azure AD authenticated users:
<https://developer.hashicorp.com/terraform/language/settings/backends/azurerm>

Logout and login again to Azure CLI to refresh the token.

Create an App Service

- In your **locals.tf** get the subnet ids for the app service and paas subnets.
- In the **app.tf** use the `azurerm_app_service_virtual_network_swift_connection` resource to connect the app service to the app service subnet:

https://registry.terraform.io/providers/hashicorp/azurerm/latest/docs/resources/app_service_virtual_network_swift_connection

Create an App Service

- Create a **pe.tf** file to define the private endpoint
- https://registry.terraform.io/providers/hashicorp/azurerm/latest/docs/resources/private_endpoint
- The subnet Id should be the Paas Subnet id
- For the private service connection, you just have to define:
 - A name
 - The private connection resource id, which is the id of your App Service
 - Sub resources name array should contains only "sites"
 - Set the manual connection boolean to false.




Create an App Service

If you succeed you should in the Networking section of your App Service these features On:


Inbound Traffic

Manage access and incoming services.

Features

 Access restriction	✓ On
 App assigned address	● N/A ⓘ
 Private endpoints	✓ On

Inbound address

10.0.0.4 

Web App

These custom domains direct traffic to your web app.



Domains

app-dev-we-prj1-hol-01.azurewebsites.net


Outbound Traffic

Set up calls to app dependencies like databases.

Features

 VNet integration	✓ On
 Hybrid connections	■ Off

Outbound addresses ⓘ

20.238.219.159,20.238.219.235,20.31.145.63,20.238.220.98 

[...Show more](#)

Terraform Modules

What is a module?

- A module is a **reusable** Terraform template
- Allow re-use of **certified** and **standardized** infrastructure as **code**
- **Can be published** at Public Module Registry at <https://registry.terraform.io/>
- **Private Repo** can be established using GitHub, Bitbucket , S3 buckets,....

Structure of a module



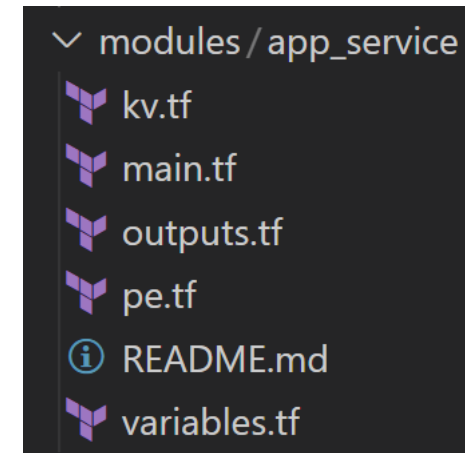
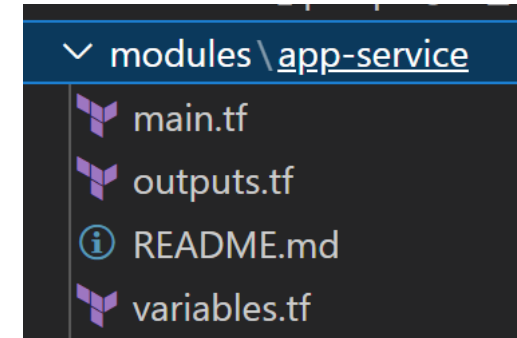
- **File:** variables.tf
- Parameters for module

- **File:** main.tf
- Resources and Data blocks

- **File:** outputs.tf
- Return results, others may reuse for their processing

Structure of a module

- If you follow the **convention**, you should have this kind of structure:
- For more clarity you can **split the main.tf** file to avoid having one big file with everything defined.
- **Avoid** using **data sources** inside a module this can create dependency resolution **issues** for terraform




Module supported sources

What	How
Local Paths	<pre>module "alert_group" { source = "./modules/alert_action_group" }</pre>
Terraform Registry	<pre>module "consul" { source = "hashicorp/consul/aws" version = "0.1.0" } module "consul" { source = "app.terraform.io/example-corp/k8s-cluster/azurerm" version = "1.1.0" }</pre>
GitHub	<pre>module "consul" { source = "app.terraform.io/example-corp/k8s-cluster/azurerm" version = "1.1.0" } module "consul" { source = "git@github.com:hashicorp/example.git" }</pre>
Generic git repository	<pre>module "vpc" { source = "git::https://example.com/vpc.git" } module "storage" { source = "git::ssh://username@example.com/storage.git" }</pre>

Pass parameters to a module

- Run a **terraform init again to install** the local module
- Example of using a module locally:



```
module "resource_name" {  
  source = "../modules/resource_folder_name"  
  # Pass the parameters based on the variables.tf file of your module  
  sku = "Basic"  
  suffix_name = "dev-app-01"  
}
```

Refactoring into a module

- Modules are **referenced in Terraform state**
- Refactoring into a modules has an **impact** on the **state**

```
Plan: 4 to add, 0 to change, 4 to destroy.
```

- Resource names can be **changed** or **refactored** into modules
- **Define a tf file** (moved.tf for instance) when extracting a module during refactoring

```
moved {  
  from = azurerm_storage_account.myStorageAccount  
  to   = module.storageAccount.azure_rm_storage_account.storeacc  
}
```

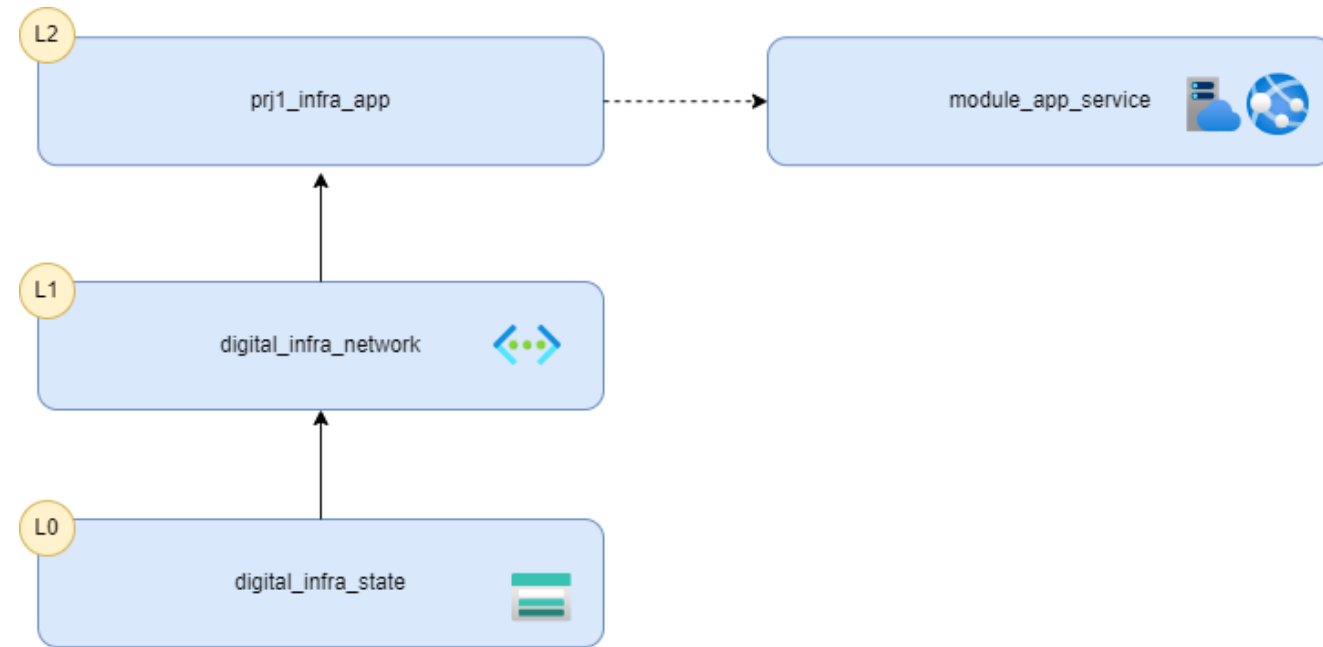
Refactoring into a module

- Map the **old resource name** to the **new one** so terraform can **update the state**
- **Use terraform plan** to detect the old new naming
- If your moved declarations are correct your **plan will not** have any resources creation or deletion
- **After** this process you can remove your .tf file with all the **moved** declarations

Hands On Lab!

Migrate the App Service to a module

- Back in your **prj1_infra_app** create a new folder called **modules** and inside another one called **app_service**
- The goal is to provide a module for app service that manage the VNet integration.
- In this module, create a **main.tf** and put the **app.tf** file content into it.
- Create a **variables.tf** and an **outputs.tf**
- Move also the **pe.tf** inside this module



Migrate the App Service to a module

- Update all the local properties in the modules
- Make the app service plan sku a parameter of your module
- Add two outputs to your module: the app service id and app service default hostname
- **Reference** the new module in the **app.tf** of your **prj1_infra_app** folder
- Make sure to migrate the module **without** creating or deleting any resource
- You may have to use a **lifecycle** block
- Run **terraform plan** and **terraform apply** you should have no changes

Provisionners

Provisioners

- Provisioners provide access to **non-terraform commands** and aren't **fully managed by it's declarative syntax or state.**

Examples:

- **local-exec** – run a command locally
- **remote-exec** – run a ssh command on a remote system
- **file** – copy files or directories to remote machine
- All provisioners support at least the following arguments:
 - **when** = destroy – execute only on destroy
 - **on_failure** = fail / continue – default is fail

local-exec

- Execute scripts or other programs on the host executing terraform.

Argument	Type	Description
command	string, required	<ul style="list-style-type: none">• Command to execute• Evaluated in a shell, and can use environment variables or Terraform variables
working_dir	string, optional	<ul style="list-style-type: none">• Working directory where the command will be executed
interpreter	string list, optional	<ul style="list-style-type: none">• List of interpreter arguments used to execute the command.• First argument is the interpreter itself, as relative path to the current working directory or as an absolute path.• The remaining arguments are appended prior to the command.• Example: ["/bin/bash", "-c", "echo foo"]• If interpreter is unspecified, sensible defaults will be chosen based on the system OS
environment	object, optional	<ul style="list-style-type: none">• Block of key value pairs representing the environment of the executed command.

local-exec

- Example: command + interpreter

```
resource "null_resource" "this" {  
  provisioner "local-exec" {  
    command = "Get-Date > completed.txt"  
    interpreter = ["PowerShell", "-Command"]  
  }  
}
```

- Example: command + environment

```
resource "null_resource" "this" {  
  provisioner "local-exec" {  
    command = "echo $FOO $BAR $BAZ >> env_vars.txt"  
  
    environment = {  
      FOO = "bar"  
      BAR = 1  
      BAZ = "true"  
    }  
  }  
}
```


remote-exec

- Execute scripts or other programs on a remote system.
- Only one of the following arguments allowed:

Argument	Type	Description
inline	string list	<ul style="list-style-type: none">• List of executed commands in specified order
script	string	<ul style="list-style-type: none">• Path (relative or absolute) of a local script• Local script will be copied to the remote resource and then executed
scripts	string list	<ul style="list-style-type: none">• List of paths (relative or absolute) of a local script• Local scripts will be copied to the remote resource and then executed in specified order

remote-exec

- Execute scripts or other programs on the remote system
- Examples:
 - Change some permissions and ownership
 - Run a script with some environment variables



```
resource "some_resource" "this" {  
  ...  
  provisioner "remote-exec" {  
    inline = [  
      "sudo chown -R ${var.admin_username}:${var.admin_username}  
/var/www/html",  
      "chmod +x *.sh",  
      "PLACEHOLDER=${var.placeholder} WIDTH=${var.width}  
HEIGHT=${var.height} PREFIX=${var.prefix} ./deploy_app.sh",  
    ]  
  }  
  ...  
}
```


file

- Copy files or directories to a remote system

Argument	Type	Description
source	String	<ul style="list-style-type: none">• Source file or directory
content	String	<ul style="list-style-type: none">• Direct content to copy on the destination• If destination is a file, the content will be written on that file.
destination	string	<ul style="list-style-type: none">• Destination path to write to on the remote system• Will be evaluated by the remote system, rather than by Terraform itself

file


- **Copies files or directories** onto the remote system
- Support for **SSH** and **WinRM connections**



```
provisioner "file" {  
  source = "files/"  
  destination = "/home/${var.admin_username}/"  
  connection {  
    type      = "ssh"  
    user      = "${var.admin_username}"  
    password  = "${var.admin_password}"  
    host      = "${azurerm_public_ip.catapp-pip.fqdn}"  
  }  
}
```


Architecture

Layered Deployment



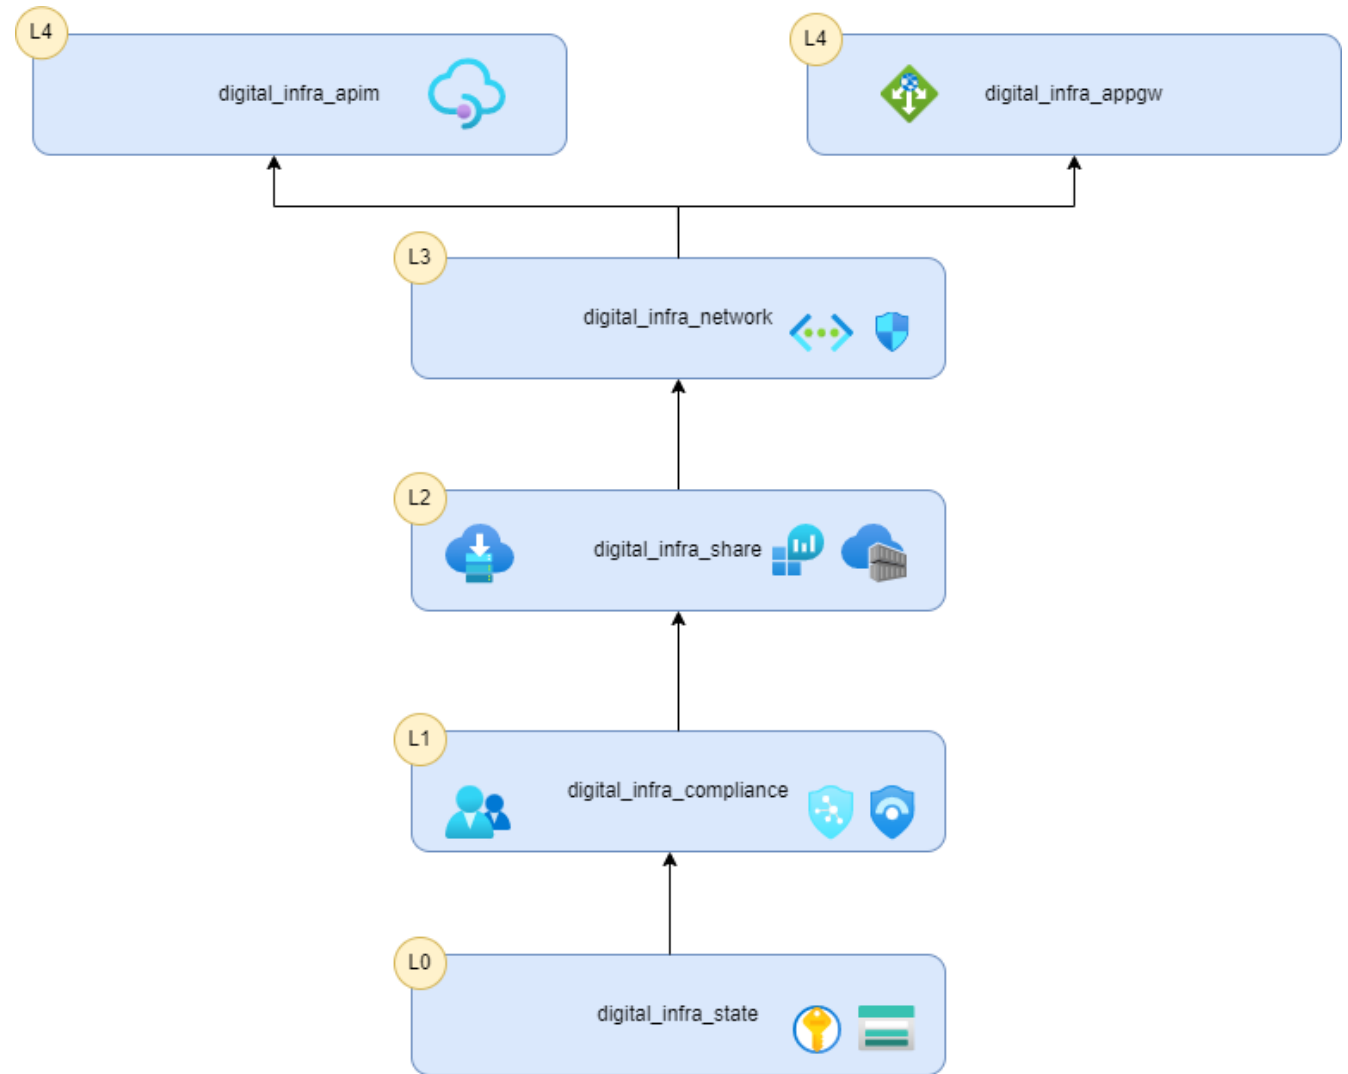
Layer	Components
0	Bootstrap with init script / DevOps Agent
1	Security & Compliance
2	Core Services & Network
3	Shared Services (Compute, Data Server)
4	Applications (DB, WebApp, Queue,...)

Layered Deployment

- Reduce the scripts to a **specific scope**
- **Speed** things up
- **Split responsibilities** and permissions between teams.
- Each level has **its own state**

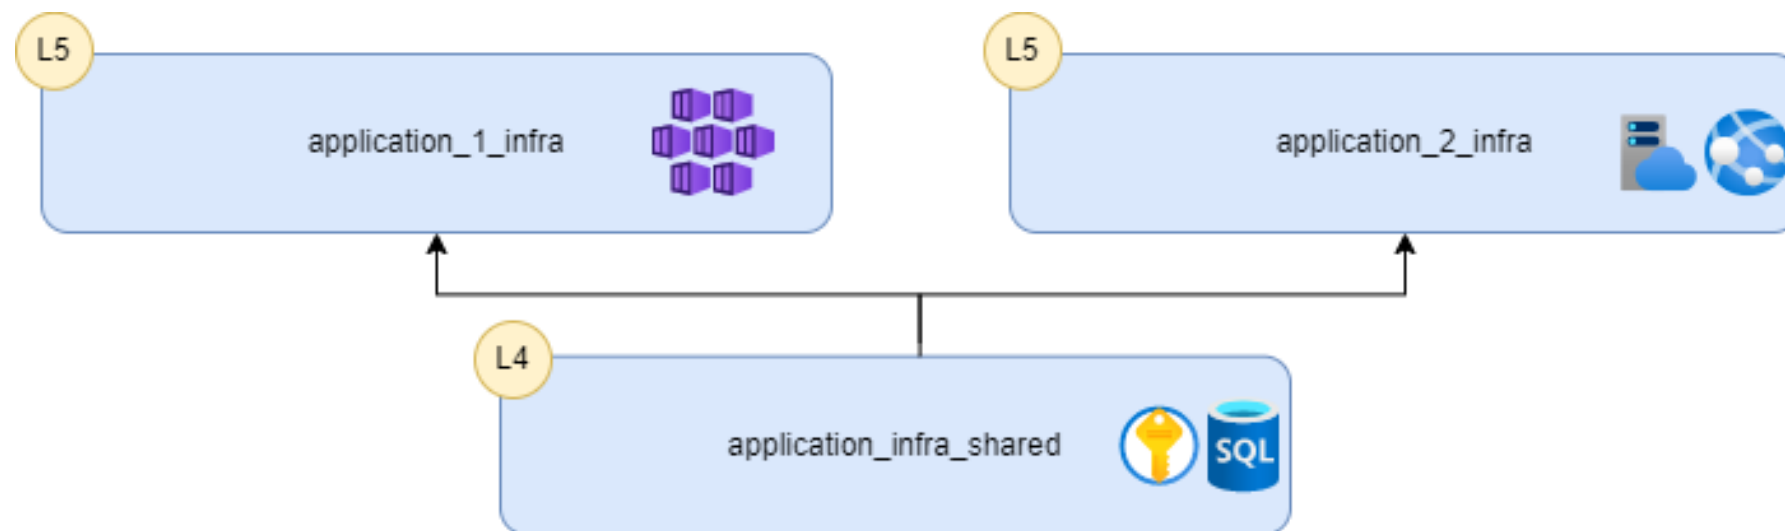
Layered Deployment: Example

- Every layer has his own tfstate **except the first one**
- Each layer has **outputs exposed** to the next layer.
- You can add **more layers** to fit your needs
- Example:
 - L1 expose outputs to L2 / L3...
 - L3 can consume outputs from L2 / L3




Layered Deployment: Example

- **Applications layers** can be splitted to share resources
- Only when **business requirements** are compatible



Layered Deployments and maturity



DevOps


Branch & deployment strategy

- Look at 10 projects you will find 10 different deployments
- The **best way** depends on **your project**
- Do a **brainstorming** before starting your project

Decide on a branch strategy

- **Main** branch is rolling out the code to **production** infrastructure
- **Develop** branch is rolling out the code to **staging** infrastructure
- **Feature** branch can be manually rolling out in a sandbox / dev subscription
- When the project grows, feature branch are rolling out using **CI/CD** in a dev environment **like the other branches.**

Branching Strategy




Decide on a deployment strategy

- **Create stages** with own environment variables
- Use **managed identity** for each environment
- **Do not define** backend configuration in your terraform code
- Create **flexible** pipelines or workflows, so a new stage for a new environment can be added **easily**
- Test your Terraform Code **before rollout** in a dev environment
- When working on **new features** be careful, when you deploy in a testing environment to not erase the deployment made by your teammate
- When the project grows, because of the **dependencies**, it becomes more and more **complicated** to work outside of a deployment environment common to the whole team

Deployment strategy

Builds (Terraform Code Freeze)

Releases (Terraform init, plan, apply)



Build strategy


- **Freeze** your Terraform (package, tag, versioning)
- **Dynamic** Terraform **plan** will be done during release
- Let you **downgrade** environments as needed
- Output your plan in a file to have something **static to deploy**

Release strategy



Terraform Rollout

- To change some resources **Terraform deletes** them and **recreates** them
- **Always check the plan** and test it on non-productive data



Many tools to start

- Azure DevOps, GitHub Actions, etc..




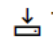


- Of course, you can reproduce the commands run by the tasks and actions directly using **basic command lines**

Azure DevOps

- Pre-build **tasks** are available to help you, install terraform and run commands in your pipelines



Terraform

Microsoft DevLabs  microsoft.com |  75,255 installs |   (27) | Free

Install terraform and run terraform commands to manage resources on Azure, AWS and GCP.

Get it free

<https://marketplace.visualstudio.com/items?itemName=ms-devlabs.custom-terraform-tasks>

GitHub Actions

- Pre-build **actions** are available to help you, install terraform and run commands in your workflows.



GitHub Action

HashiCorp - Setup Terraform

v2.0.3

Latest version

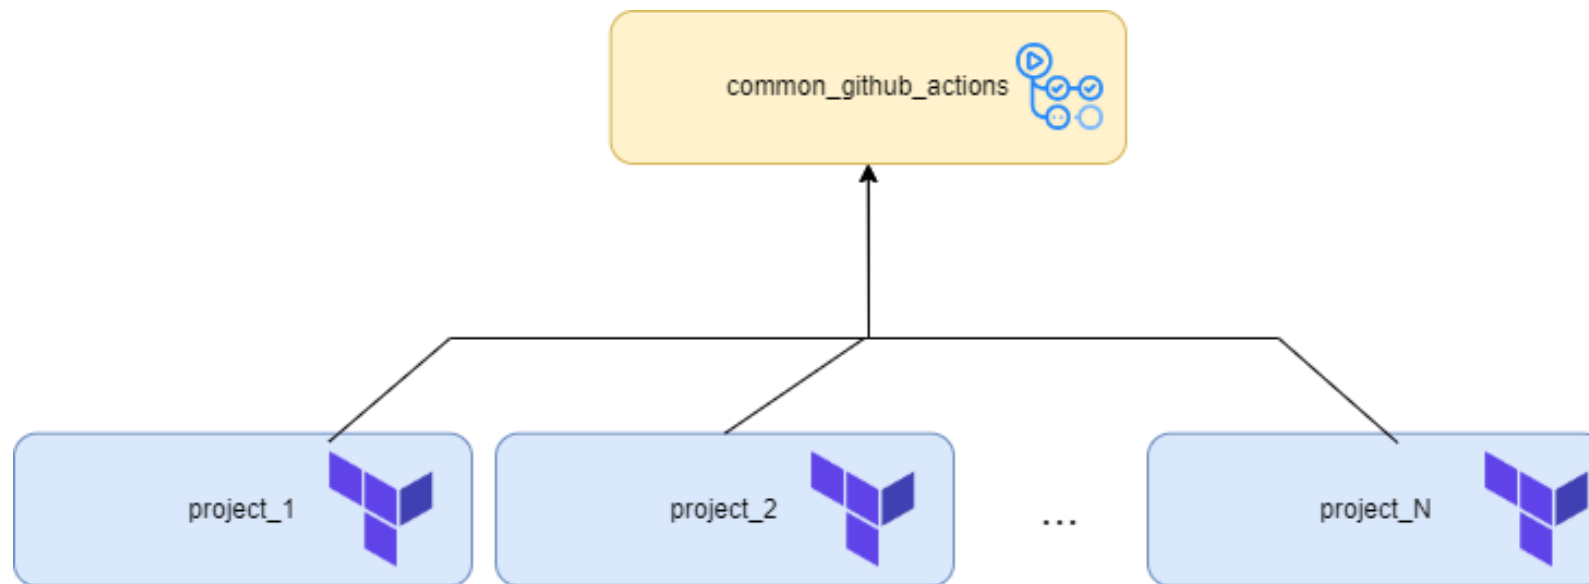
Use latest version



<https://github.com/marketplace/actions/hashicorp-setup-terraform>

Mutualize your GitHub Workflows

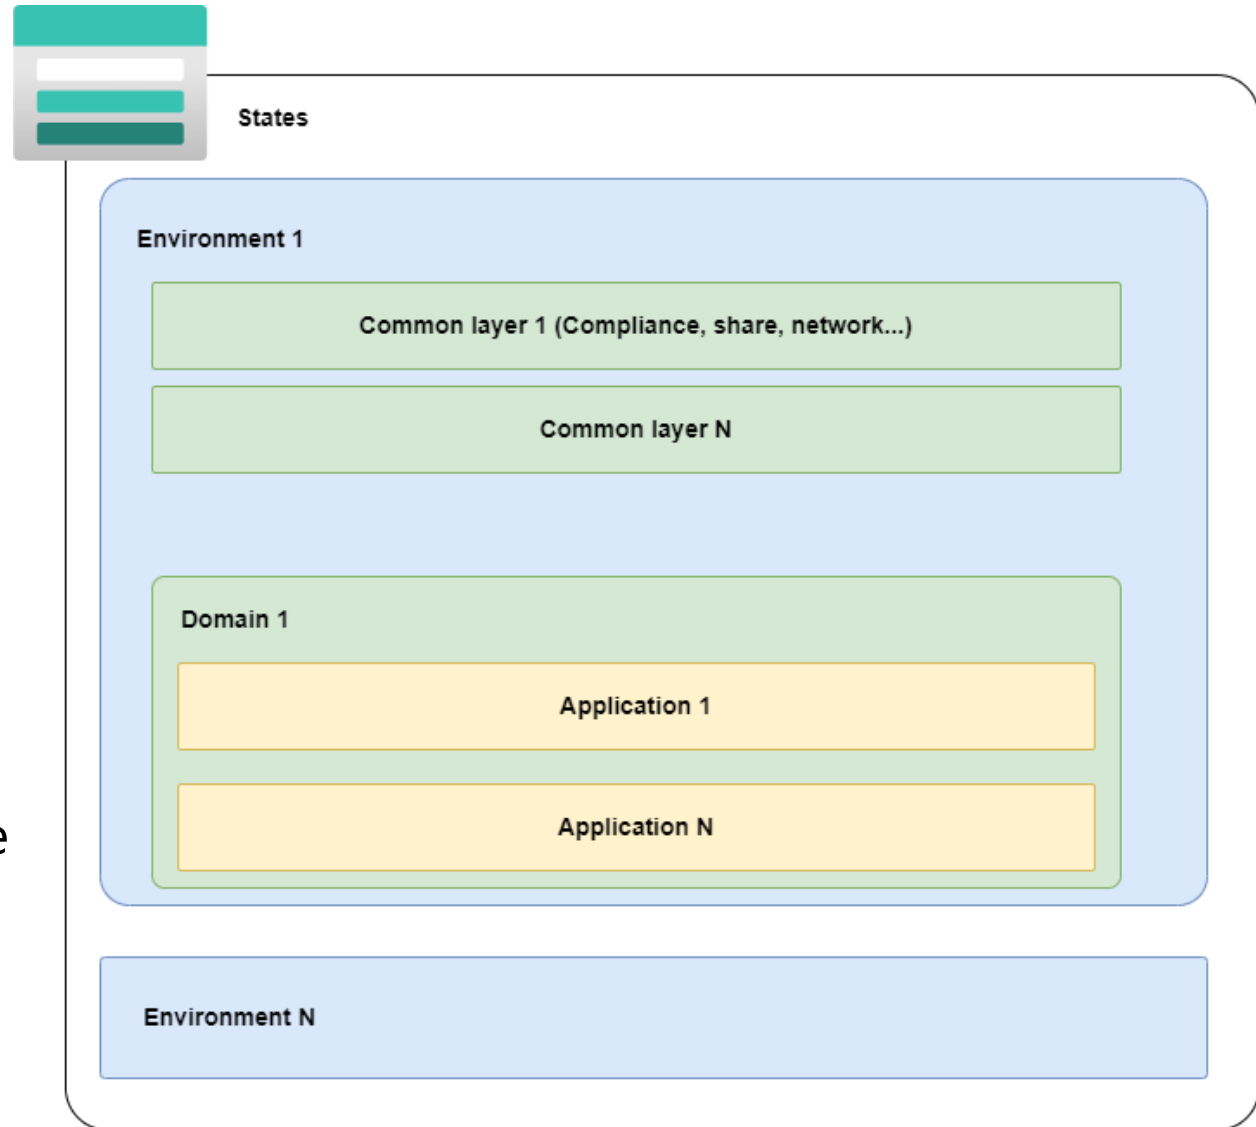
- The terraform workflows is always the same, try to mutualize it



Structure the storage of your states

- Split your states **logically**
- Create a **structure convention** to store your states
- **Automatically predict** the path of your state
- For instance, for an application 1 the path to the state in the Storage Account can be:

`/states/dev/domain1/application1/terraform.state`



Hands On Lab!

Migrate to GitHub Actions

- Choose one project for instance the **prj1_infra_app** and add a .gitignore to it: <https://www.toptal.com/developers/gitignore/api/terraform>
- **Add any other files** like plan or .terraform* that are missing to the .gitignore if necessary
- Upload one project to GitHub
- Create a github **workflow** for the project
- Create a **job** to run a Terraform init, validate and plan
- Create another job to run a Terraform **apply** if changes where detected
- Use **terraform plan --detailed-exitcode**
- Use the **identity method** you prefer

Tools

Terraform-docs



- Generate documentation from Terraform modules in various output format
- <https://github.com/terraform-docs/terraform-docs/>
- One config file (.terraform-docs.yml) is used to set the template and the output of the documentation
- One simple command line to generate the documentation:



```
terraform-docs -c .terraform-docs.yml .
```


TFLint

- [terraform-linters/tflint: A Pluggable Terraform Linter \(github.com\)](https://github.com/terraform-linters/tflint)
- TFLint is a framework, and each feature is provided by plugins:
 - Find possible errors, like illegal instance types
 - Warn about deprecated syntax, unused declarations
 - Enforce best practices, naming conventions
 - Plugins for the Azure Rm provider is available: <https://github.com/terraform-linters/tflint-ruleset-azurerm>
- Use it **after** running terraform validate command
- A tflint Action is also available: <https://github.com/terraform-linters/setup-tflint>

TFSec

- Made by Aqua Security
- <https://aquasecurity.github.io/tfsec/>
- Static analysis security scanner



Checkov

- Made by bridgecrew (Prima Cloud)
- <https://www.checkov.io/>
- Checkov is a static code analysis tool
- Includes predefined policies to check for common misconfiguration issues
- Supports the creation and contribution of custom policies.
- Integrated with GitHub Actions: <https://www.checkov.io/4.Integrations/GitHub%20Actions.html>

Terratest: End-2-end tests

- [Terratest | Automated tests for your infrastructure code. \(gruntwork.io\)](https://gruntwork.io/terratest/)
- [Test Terraform modules in Azure using Terratest | Microsoft Docs](https://docs.microsoft.com/en-us/azure/terraform/terraform-test-terratest)
- Tests are made using Go



Write test code using Go

Create a file ending in `_test.go` and run tests with the `go test` command. E.g., `go test my_test.go`.



Use Terratest to deploy infrastructure

Use Terratest to execute your real IaC tools (e.g., Terraform, Packer, etc.) to deploy real infrastructure (e.g., servers) in a real environment (e.g., AWS).



Validate infrastructure with Terratest

Use the tools built into Terratest to validate that the infrastructure works correctly in that environment by making HTTP requests, API calls, SSH connections, etc.



Undeploy

Undeploy everything at the end of the test.

Infracost

- Infracost shows cloud cost estimates for Terraform.
- It lets engineers understand costs **before making changes**
 - In the **terminal**
 - In **VS Code**
 - In **pull requests**
- <https://www.infracost.io/docs/>



Infracost

Infracost  infracost.io |  12,738 installs |  (5) | Free

Cloud cost estimates for Terraform in your editor

[Install](#) [Trouble Installing?](#) 

Hands On Lab!

Play with terraform-docs, TFLint and TFSec

- Use terraform-docs: <https://terraform-docs.io/user-guide/introduction/>
- Add a **.terraform-docs.yml** configuration file to at least one project.
- Try to generate a basic documentation in a file called **DOCS.md**
- Use TFLint: <https://github.com/terraform-linters/tflint>
- Add a **.tflint.hcl** configuration file to at least one project.
- Add the **azurerm** plugin
- Try to disable a **tflint rule**
- Use TFSec: <https://github.com/aquasecurity/tfsec>
- Run TFSec to at least one project

Migrate to Terraform

Migrate to Terraform

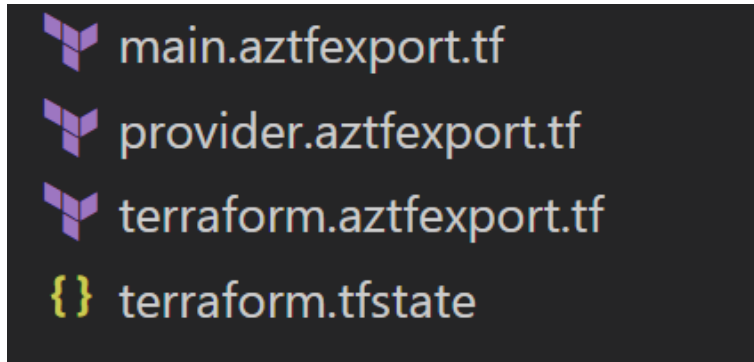
- There is **no perfect tools** to do this
- Multiple technics can be used:
- **Manually transform** the properties from ARM files to Terraform files
- **Manually transform** the properties from Bicep files to Terraform files
- Use **Microsoft Azure Export for Terraform** to speed up the process
- Use Terraform **import** command release with the version 1.5+
<https://www.hashicorp.com/blog/terraform-1-5-brings-config-driven-import-and-checks>

Microsoft Azure Export for Terraform

- <https://github.com/Azure/aztfexport>
- This will do the **big part** of the job but needs to be rework after
- It **extracts the configuration** from Azure and generate a Terraform code
- The Terraform configurations generated by **aztfexport** are **not meant to be comprehensive**
- It do **not ensure** that the infrastructure can be fully reproduced from the generated configurations.

Microsoft Azure Export for Terraform

- This will save some time, but **you must double check** what was generated
- **Split it** correctly in multiple files and eventually in layers.
- This will generate 4 files:



```
main.aztfexport.tf  
provider.aztfexport.tf  
terraform.aztfexport.tf  
{ } terraform.tfstate
```

- All the infrastructure imported from Azure will be in one file: **main.aztfexport**

Terraform import


- This command available since Terraform 1.5
- This generate all the code based on the resources you target
- Some default properties are generated but it can be a good starting point



```
terraform plan -generate-config-out=generated.tf
```

Hands On Lab!

Migrate resource to Terraform



- Let's simulate the migration of a resource: Create an Azure Container Registry with Basic SKU in a resource group
- Then, try **one** of these options to do the Terraform code:

Migrate ARM or Bicep template manually extracted to Terraform

Try to generate it with **aztfexport**
<https://github.com/Azure/aztfexport>

Use Terraform **import** to generate it
<https://developer.hashicorp.com/terraform/tutorial/s/state/state-import#define-import-block>

- Create a new folder called **digital_infra_share** and add this terraform code to it.

Thank you.